

A rapid and efficient learning rule for biological neural circuits

Eren Sezener^{*1}, Agnieszka Grabska-Barwińska^{*1}, Dimitar Kostadinov^{*2},
Maxime Beau², Sanjukta Krishnagopal³, David Budden¹,
Marcus Hutter¹, Joel Veness¹, Matthew Botvinick^{1,3},
Claudia Clopath^{†4}, Michael Häusser^{†2}, Peter E. Latham^{†3}

^{*}†Equal contribution

¹DeepMind

²Wolfson Institute for Biomedical Research, University College London

³Gatsby Computational Neuroscience Unit, University College London

⁴Department of Bioengineering, Imperial College London

August 17, 2022

Abstract

The dominant view in neuroscience is that changes in synaptic weights underlie learning. It is unclear, however, how the brain is able to determine which synapses should change, and by how much. This uncertainty stands in sharp contrast to deep learning, where changes in weights are explicitly engineered to optimize performance. However, the main tool for that, backpropagation, has two problems. One is neuroscience related: it is not biologically plausible. The other is inherent: networks trained with this rule tend to forget old tasks when learning new ones. Here we introduce the Dendritic Gated Network (DGN), a variant of the Gated Linear Network, which offers a biologically plausible alternative to backpropagation. DGNs combine dendritic ‘gating’ (whereby interneurons target dendrites to shape neuronal responses) with local learning rules to yield provably efficient performance. They are significantly more data efficient than conventional artificial networks, and are highly resistant to forgetting. Consequently, they perform well on a variety of tasks, in some cases better than backpropagation. Importantly, DGNs have structural and functional similarities to the cerebellum, a link that we strengthen by using *in vivo* two-photon calcium imaging to show that single interneurons suppress activity in individual dendritic branches of Purkinje cells, a key feature of the model. Thus, DGNs leverage targeted dendritic inhibition and local learning – two features ubiquitous in the brain – to achieve fast and efficient learning.

31 1 Introduction

32 A hallmark of intelligent systems is their ability to learn. Humans, for instance, are
33 capable of amazing feats – language acquisition and abstract reasoning being the most
34 notable – and even fruit flies can learn simple reward associations [1, 2]. It is widely
35 believed that this learning is implemented via synaptic plasticity. But which synapses
36 should change in response to, say the appearance of a reward, and by how much?
37 This is especially hard to answer in humans, who have about 10^{14} synapses, but it is
38 hard even in fruit flies, which have about 10^7 – corresponding to 10 million adjustable
39 parameters.

40 One answer to this question is known: introduce a loss function (a function that
41 measures some aspect of performance, with higher performance corresponding to lower
42 loss), compute the gradient of the loss with respect to the weights (find the direction
43 in weight space that yields the largest improvement in performance), and change the
44 weights in that direction. If the weight changes are not too large, this will, on average,
45 reduce the loss, and so improve overall performance.

46 This approach has been amazingly successful in artificial neural networks, and has
47 in fact driven the deep learning revolution [3]. However, the algorithm for computing
48 the gradient in deep networks is not directly applicable to biological systems, as first
49 pointed out by [4, 5] (see also recent reviews [6–8]). There are several reasons for this.
50 First, to implement backpropagation [9–11], referred to simply as backprop, neurons
51 would need to know their outgoing weights. Second, backprop requires two stages:
52 a forward pass (for computation) and a backward pass (for learning). Moreover, in
53 the backward pass an error signal must propagate from higher to lower areas, layer
54 by layer (Fig. 1A), and during that backward pass information from the forward pass
55 must remain in the neurons. However, biological neurons do not know their outgoing
56 weights, and there is no evidence for a complicated, time-separated backward pass.

57 Backprop also leads to another problem, at least in standard deep learning setups:
58 it adapts to the data it has seen most recently, so when learning a new task it forgets
59 old ones [12]. This is known as catastrophic forgetting, and prevents networks trained
60 with backprop to display the lifelong learning that comes so easily to essentially all
61 organisms [13, 14].

62 Driven in part by the biological implausibility of backprop, there have been several
63 proposals for architectures and learning rules that might be relevant to the brain.
64 These include feedback alignment [15, 16], creative use of dendrites [17, 18], multiplexing
65 [19], and methods in which the error signal is fed directly to each layer rather than
66 propagating backwards from the output layer [20–28]. A particularly promising method
67 that falls into the latter category is embodied in Gated Linear Networks [29, 30]. These
68 networks, which were motivated from a machine learning rather than a neuroscience
69 perspective, have obtained state-of-the-art results in regression and denoising [31],
70 contextual bandit optimization [32], and transfer learning [33].

71 In Gated Linear Networks (GLNs), the goal of every neuron, irrespective of its layer,
72 is to predict the target output based on the input from the layer directly below it. This
73 is very different from backprop, in which neurons in intermediate layers extract features
74 that make it easier for subsequent layers to predict the target (compare Figs. 1A and B).

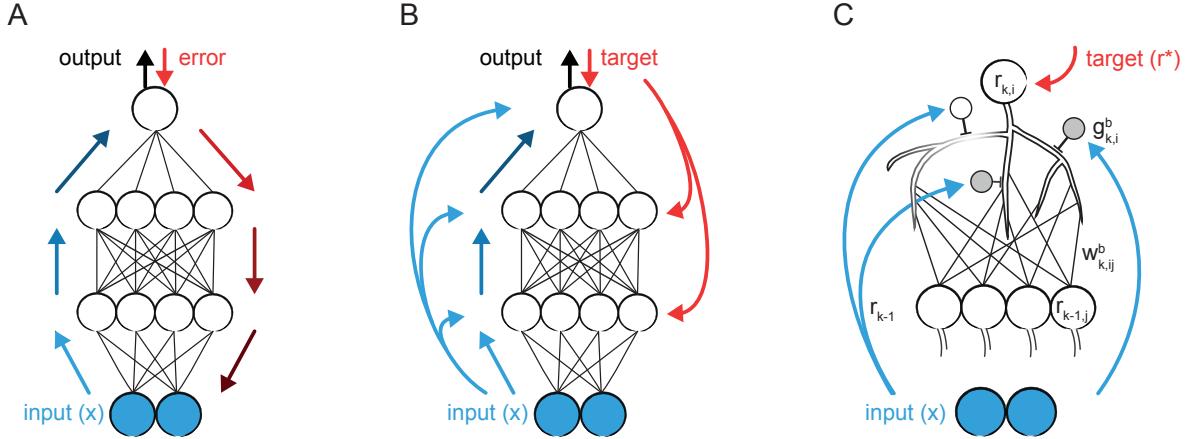


Figure 1: Comparison of multi-layer perceptrons (MLPs) and Dendritic Gated Networks (DGNs). In all panels the blue filled circles at the bottom correspond to the input. **A.** MLP. Blue arrows show feedforward computations; red arrows show the error propagating back down. **B.** DGN. As with MLPs, information propagates up, as again shown by the blue arrows. However, rather than the error propagating down, each layer receives the target output, which it uses for learning. Connections from the input to each layer (light blue arrows) support the gating. **C.** A single postsynaptic neuron in layer k of a DGN, along with several presynaptic neurons in layer $k - 1$. Each branch gets input from all the presynaptic neurons (although this is not necessary), and those branches are gated on and off by inhibitory interneurons which receive external input. The white interneuron is active, so its corresponding branch is gated off, as indicated by the light gray branches; the gray neurons are not active, so their branches are gated on.

Gated Linear Networks are thus particularly suitable for biologically plausible learning: every neuron is essentially part of a shallow network, with no hidden layers, for which the delta rule [34] – a rule that depends only on presynaptic and postsynaptic activity – is sufficient to learn.

To implement these local learning rules, the target activity (a scalar) is sent to every neuron, in every layer of the network (Fig. 1B, red arrows). This is typical of a large class of learning rules [20, 21, 23–27]. Completely atypical, though, is the role of the external input. It is used for gating the weights: each neuron has a bank of weights at its disposal, and the external input determines which one from that bank is used. For example, a neuron might use one set of weights when the visual input contains motion cues predominantly to the right; another set of weights when it contains motion cues predominantly to the left; and yet another when there are no motion cues at all. (Note that this example is over-simplified: in practice the input is high dimensional, and the mapping from external input to the chosen set of weights contains very little structure; see Fig. 2C.)

Having a “look-up” table, in which each input corresponds to a particular set of weights, is inconsistent with what we see in the brain. However, we can attain the performance of Gated Linear networks by gating dendritic branches on and off, using

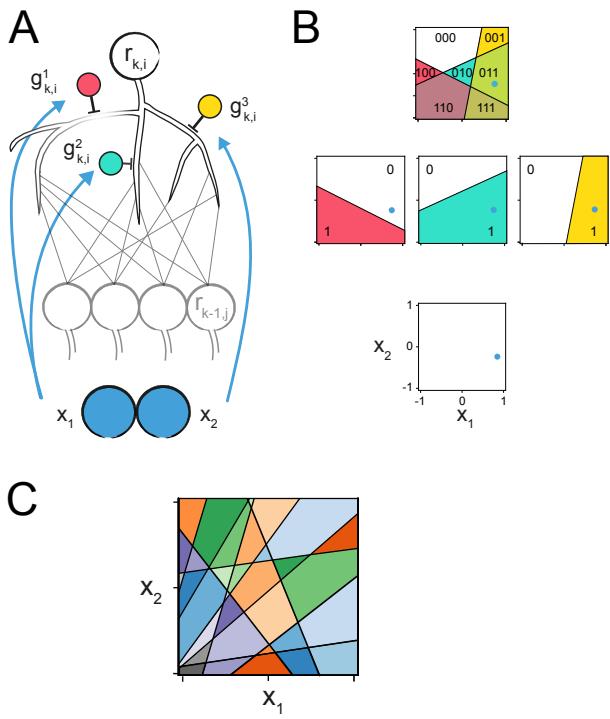


Figure 2: Random mapping (here implemented with so-called half-space gating, which we use throughout the paper; see Eq. (3)), shown for two dimensional input for clarity. (In realistic cases the input is, of course, high dimensional.) **A.** Two input neurons (blue) connect to three dendritic branches, and gate them either on or off. Each gate divides the two dimensional input into two half-spaces, one of which shuts down the gate, silencing the corresponding dendritic branch. **B.** For the input on this particular trial (blue dot in bottom square), the gate $g_{k,i}^1$ (red) is off, while $g_{k,i}^2$ (teal) and $g_{k,i}^3$ (yellow) are on, as indicated by the three squares with the corresponding colors. The top square shows a summary of the possible combinations of weights. Each of the seven regions has a different combination, making it possible to implement a large range of input-output mappings. **C.** A more realistic case of 10 branches. In DGNs, each coloured region corresponds to a linear combination of 10 sets of weights. This is in contrast to GLNs, which use a separate set of weights for each region.

inhibitory neurons, in an input-dependent manner (Figs 1C and 2). We thus replace the weight look-up mechanism of GLNs with linearly additive dendritic weights, and refer to these networks as Dendritic Gated Networks (DGNs). Perhaps surprisingly, the mapping from the input to the dendritic branches is completely random, so the input isn't chosen to target specific branches (see, for example, Fig. 2C). This stands in sharp contrast to architectures where the gating is learned [35]. But the unlearned, random mapping is in fact a key ingredient, as it allows DGNs to represent essentially arbitrary nonlinear functions efficiently. Moreover, this gating makes DGNs especially resistant to forgetting. In particular, when data comes in separate "tasks", DGNs can learn new ones without forgetting the old. Finally, the loss is a convex function of the weights for each unit (see Supplementary Information, "Convexity"), as it is in Gated Linear Networks [29]. Convexity is an extremely useful feature, as it enables DGNs, like the Gated Linear Networks on which they are based, to learn quickly.

Below we describe multi-layer Dendritic Gated Networks in detail – both the architecture and the learning rule. We then train them on four tasks: two on which vanilla feedforward networks trained with backprop typically exhibit catastrophic forgetting, and two relevant to the cerebellum. We map the proposed learning rule and the associated architecture to the cerebellum because 1) the climbing fibers provide a well-defined feedback signal; 2) its input-output function is relatively linear [36–38]; and 3) molecular layer interneurons could act as gates [39–48]. Finally, we demonstrate experimentally that a key prediction of the DGN – suppression of individual dendritic branches by interneurons – is observed in cerebellar Purkinje cell spiny branchlets *in vivo*. Thus, our theoretical and experimental results draw a specific link between learning in DGNs and the functional architecture of the cerebellum. The generality of the DGN architecture should also allow this algorithm to be implemented in a range of networks in the mammalian brain, including the neocortex.

2 Results

2.1 Dendritic Gated Networks

Dendritic Gated Networks, like conventional deep networks, are made up of multiple layers, with the input to each layer consisting of a linear combination of the activity in the previous layer. Unlike conventional deep networks, however, the weights are controlled by external input, via gating functions, denoted $g(\mathbf{x})$; those functions are implemented via dendritic branches (Figs. 1B, C and 2). This results in the following network equations. The activity (i.e., the instantaneous firing rate) of the i^{th} neuron in layer k , denoted $r_{k,i}$, is

$$r_{k,i} = \phi \left(\sum_{b=1}^{B_{k,i}} g_{k,i}^b(\mathbf{x}) \sum_{j=0}^{n_{k-1}} w_{k,ij}^b h_{k-1,j} \right), \quad (1)$$

with the synaptic drive, $h_{k-1,j}$, given in terms of $r_{k-1,j}$ as

$$h_{k-1,j} = \phi^{-1}(r_{k-1,j}). \quad (2)$$

121 Here $\phi(\cdot)$ is the activation function (either identity or sigmoid), $r_{k,i}$ is the activity of
 122 i^{th} neuron in layer k (with $r_{k,0}$ set to 1 to allow for a bias), $B_{k,i}$ is the number of
 123 branches of neuron i in layer k , $w_{k,ij}^b$ is the weight from neuron j in layer $k-1$ to
 124 the b^{th} branch of neuron i in layer k , n_k is the number of neurons in layer k , and
 125 $g_{k,i}^b(\mathbf{x})$ is the binary gating variable; depending on the external input, \mathbf{x} (taken to be
 126 an n -dimensional vector), it's either 1 (in which case the b^{th} branch of the i^{th} neuron
 127 is gated on) or 0 (in which case it's gated off). There are K layers, so k runs from 1 to
 128 K . The input to the bottom layer is \mathbf{x} – the same as the input to the gating variable.

The mapping from the input, \mathbf{x} , to the gating variable, $g_{k,i}^b(\mathbf{x})$, is not learned;
 instead, it is pre-specified, and does not change with time. In all of our simulations we
 use random half-space gating [29]; that is,

$$g_{k,i}^b(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{v}_{k,i}^b \cdot \mathbf{x} \geq \theta_{k,i}^b \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

129 where $\mathbf{v}_{k,i}^b$ and $\theta_{k,i}^b$ are sampled randomly and kept fixed throughout learning (see
 130 Methods), and “ \cdot ” is the standard dot product.

131 Note that a Dendritic Gated Network with one branch reduces to a Gated Linear
 132 Network. With a caveat: in the original formulation [29], Gated Linear Networks had
 133 a nonzero weight for each input, \mathbf{x} , which is not the case if weights are completely
 134 gated off for one of the half spaces (because in that case the weights are zero). For
 135 a detailed description of the difference between GLNs and DGNs, see Supplementary
 136 Information, “Difference between GLNs and DGNs”.

In Dendritic Gated Networks, the goal of each neuron is to predict the target output,
 denoted r^* (which is a function of \mathbf{x} ; we suppress the \mathbf{x} -dependence to reduce clutter).
 To do that, the weights, $w_{k,ij}^b$, are modified to reduce the loss, $\ell_k(r^*, r_{k,i})$. For weight
 updates we use gradient descent,

$$\Delta w_{k,ij}^b = -\eta \frac{\partial \ell(r^*, r_{k,i})}{\partial w_{k,ij}^b} \quad (4)$$

where η is the learning rate, and the updates are performed after each sample. The form
 of the loss can influence both the speed of learning and the asymptotic performance,
 but conceptually we should just think of it as some distance between r^* and $r_{k,i}$. In
 the simplest case, which is suitable for regression, ϕ is the identity ($r_{k,i} = h_{k,i}$) and the
 loss is quadratic,

$$\ell(r^*, r_{k,i}) = \frac{1}{2}(r^* - r_{k,i})^2, \quad (5)$$

so the update rule is

$$\Delta w_{k,ij}^b = \eta g_{k,i}^b(\mathbf{x})(r^* - r_{k,i}) h_{k-1,j}. \quad (6)$$

This has the form of a gated version of the delta rule [34]. For classification, a different
 loss function is more appropriate. However, the update rule still has the form of a
 gated version of the delta rule; see Methods, Sec. 4.1, for details.

140 **2.2 Simulations**

141 Equations (1) and (3) for the network dynamics and Eq. (4) for learning constitute
142 a complete description of our model. For a given problem, we just need to choose a
143 target input-output relationship (a mapping from \mathbf{x} to r^*) and specify the loss func-
144 tions, $\ell(r^*, r_{k,i})$. Here we consider four tasks. The first two, designed to illustrate the
145 resistance of DGNs to catastrophic forgetting, are classification tasks, for which we
146 use a sigmoid activation and cross-entropy loss (Methods, Sec. 4.1); the second two,
147 which are relevant to the cerebellum, are regression tasks, for which we use an identity
148 activation and quadratic loss, as just described.

149 **DGNs can mitigate catastrophic forgetting.**

150 Animals are able to acquire new skills throughout life, seemingly without compromising
151 their ability to solve previously learned tasks [13, 14]. Standard networks do not share
152 this ability: when trained on two tasks in a row, they tend to forget the first one.
153 This phenomenon, known as “catastrophic forgetting”, is an old problem [49–51], and
154 many algorithms have been developed to address it. These typically fall into two
155 categories. The first involves replaying tasks previously seen during training [51–53].
156 The second involves explicitly maintaining additional sets of model parameters related
157 to previously learned tasks. Examples include freezing a subset of weights [54, 55],
158 dynamically adjusting learning rates [56], and augmenting the loss with regularization
159 terms with respect to past parameters [57–59]. A limitation of these approaches (aside
160 from additional algorithmic and computational complexity) is that they require task
161 boundaries to be provided or accurately inferred.

162 Unlike contemporary neural networks, the DGN architecture and learning rule is
163 naturally robust to catastrophic forgetting without any modifications or knowledge of
164 task boundaries (something that has been shown for Gated Linear Networks as well
165 [30]). In Fig. 3 we illustrate, on a simple task, the mechanism behind this robustness,
166 and show how it differs from a standard multi-layer perceptron; details are given in
167 the caption.

168 To demonstrate robustness to catastrophic forgetting on a more challenging task,
169 we train a DGN on the pixel-permuted MNIST continual learning benchmark [57, 60].
170 In this benchmark, the network has to learn random permutations of the input pixels,
171 with the random permutation changing every 60,000 trials (see Methods for additional
172 details). We compare the DGN to a multi layer perceptron (MLP) with and without
173 elastic weight consolidation (EWC) [57], the latter a highly-effective method explicitly
174 designed to prevent catastrophic forgetting by storing parameters of previously seen
175 tasks. Although elastic weight consolidation is effective, it requires a very complicated
176 architecture. In addition, it must be supplied with task boundaries, so it receives more
177 information than the DGN.

178 Because MNIST has 10 digits, we train 10 different DGNs. This could be reduced
179 to 4 networks (in general \log_2 of the number of outputs) by using a more efficient code
180 – one in which each network divides the 10 digits into two classes. Alternatively, we
181 could use a single DGN where each unit has a 10 dimensional output corresponding to
182 the class probabilities. However, this is not biologically plausible, so we did not use it.

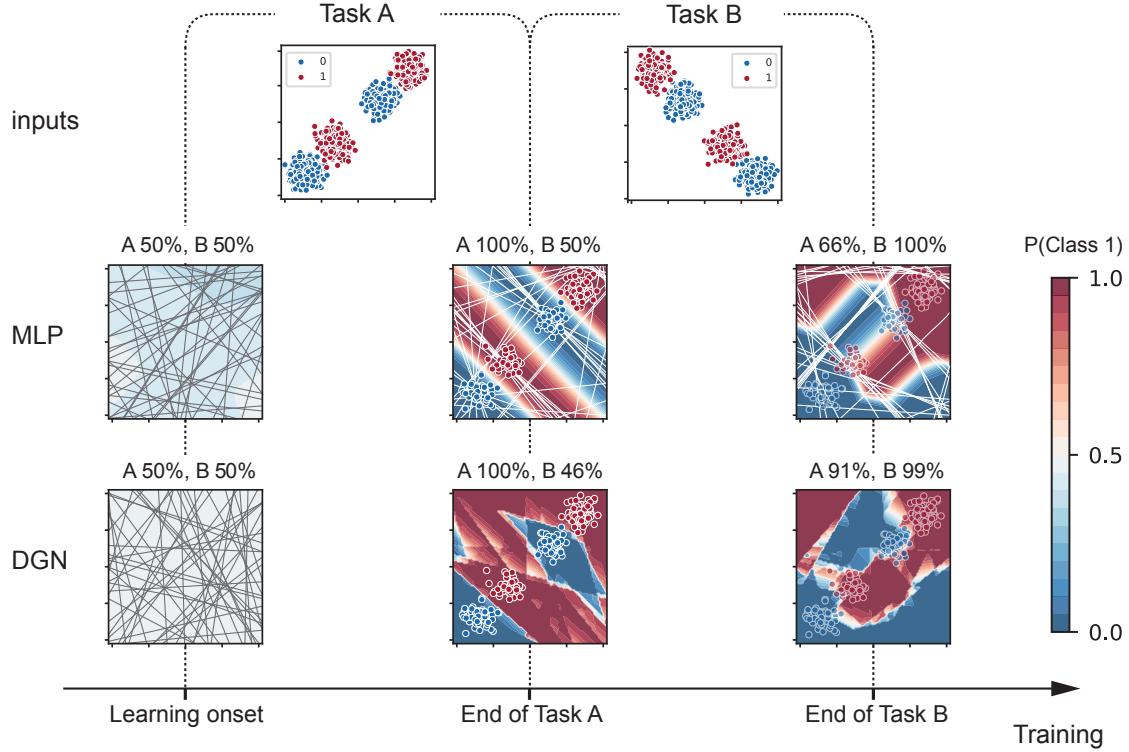


Figure 3: Comparison of the DGN to a standard multi-layer perceptron (MLP) trained with backprop. Each point on the square has to be classified as “blue” (class 0) or “red” (class 1). We consider a scenario common in the real world (but difficult for standard networks): the data comes in two separate tasks, as shown in the first row. We trained a 2-layer MLP (second row) and a 2-layer DGN (third row) on the two tasks. The output of the network is the probability of each class, as indicated by color; the percentages report the accuracy for each of the tasks. The MLP uses ReLU activation functions, so each neuron has an effective gating; the boundaries of those gates are shown in gray. The boundaries move with learning, and are plotted at the end of training of each of the tasks (white lines). The boundaries of the DGN do not move, so we plot them only in the first column. After training on Task A, most of the boundaries in the MLP are aligned at -45 degrees, parallel to the decision boundaries, which allows the network to perfectly separate the two classes. In the DGN, the boundaries do not change, but the network also perfectly separates the two classes. However, after training on Task B, the DGN retains high performance on Task A (91%), while the MLP’s performance drops to 66%. That’s because many of the boundaries changed to the orthogonal direction (45 degrees). For the DGN, on the other hand, changes to the network were much more local, allowing it to retain the memory of the old task (see samples from Task A overlaid on all panels) while accommodating the new one. The MLP has 50 neurons in the hidden layer; the DGN has 5 neurons, each with 10 dendritic branches, in the hidden layer.

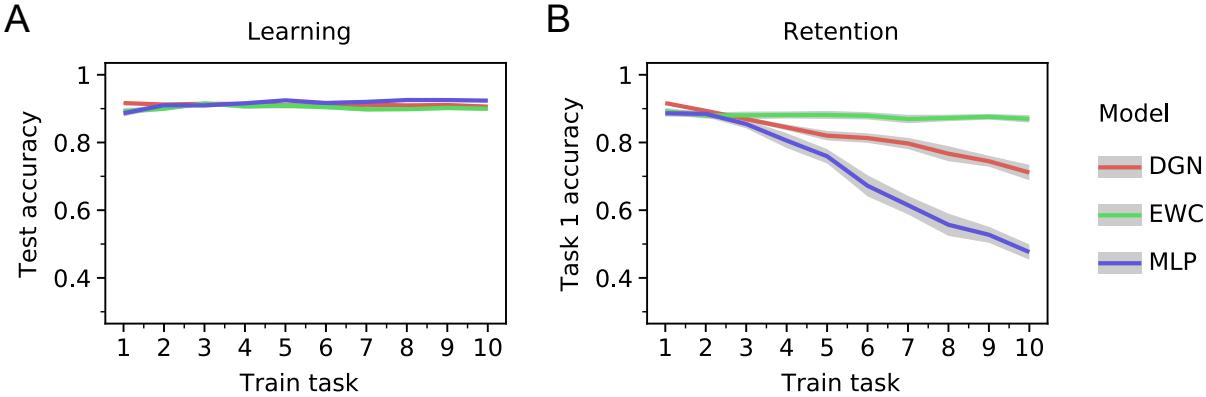


Figure 4: Learning and retention on the permuted MNIST task. The tasks are learned sequentially in a continual learning setup. **A.** Performance (on test data) for each of the 10 tasks, where a “task” corresponds to a random permutation of the pixels. **B.** Performance on the first task after each of nine new tasks is learned. As discussed in the main text, the MLP is especially bad at this task. The EWC is much better, to a large extent because it received extra information: the task boundaries. Even though the DGN was not given that information, it forgets a factor of two more slowly than the MLP. Error bars in both plots denote 95% confidence intervals over 20 random seeds.

183 Each of the 10 networks contains 3 layers, with 100, 20, and 1 neuron per layer, and
 184 there are 10 dendritic branches per neuron. The targets are categorical (1 if the digit
 185 is present, 0 if it is not), so we use binary cross-entropy rather than quadratic loss (see
 186 Methods, Sec. 4.1). We use 1000, 200, and 10 neurons per layer for the MLP (so that
 187 the number of weights match, approximately, the number weights in the DGN), with
 188 categorical cross entropy loss, both with and without elastic weight consolidation, and
 189 optimize the learning rates separately for each network.

190 Figure 4 shows the learning and retention performance of the DGN, with the MLP
 191 and EWC networks included primarily as benchmarks (neither is biologically plausible).
 192 In Fig. 4A we plot performance on each task for the three networks; as can be seen,
 193 performance is virtually identical. In Fig. 4B we investigate resistance to forgetting,
 194 by plotting the performance on the first task as the nine subsequent tasks are learned.
 195 The EWC network retains its original performance almost perfectly, the MLP forgets
 196 rapidly, and the DGN is in-between. It is not surprising that the EWC does well, as
 197 it was tailored to this task, and in particular it was explicitly given task boundaries.
 198 Somewhat more surprising is the performance of the DGN, which had none of these
 199 advantages but still forgets much more slowly than the MLP. The DGN also learns new
 200 tasks more rapidly than either the EWC or MLP networks (Supplementary Figure S3),
 201 because the loss is convex and learning is local.

202 Mapping DGNs to the Cerebellum

203 For the next two simulations we consider computations that can be mapped onto
 204 cerebellar circuitry. We focus on the cerebellum for several reasons: it is highly ex-

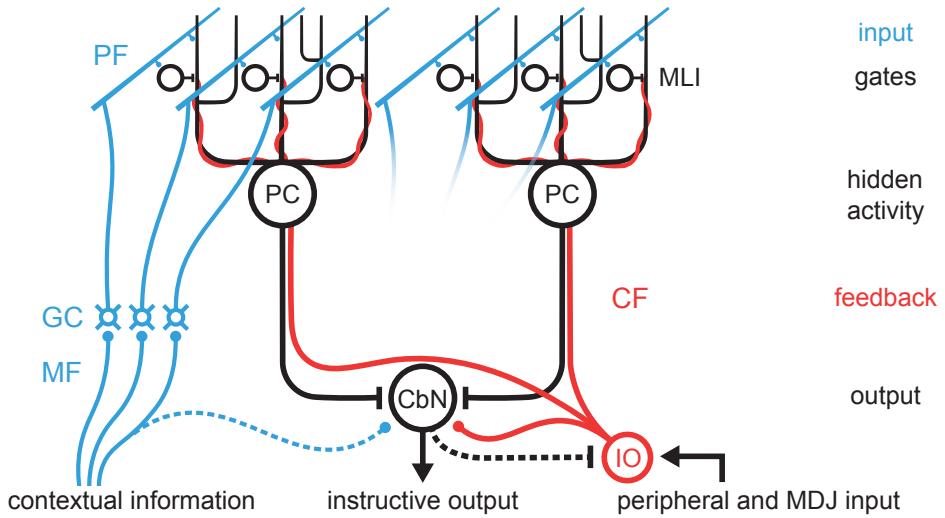


Figure 5: The cerebellum as a two layer DGN. Contextual information from the mossy fiber/granule cell (MF/GC) pathway is conveyed as input to the network via parallel fibers (PFs) that form synapses onto both the dendritic branches of Purkinje cells and molecular layer interneurons (MLIs). The inhibitory MLIs act as input-dependent gates of Purkinje cell dendritic branches. Purkinje cells converge onto the cerebellar nuclear neurons (CbNs), which constitute the output of the cerebellar network. The climbing fibers (CFs, red) originating in the inferior olive (IO) convey the feedback signal that is used to tune both the Purkinje cells, based on which inputs are gated on or off, and also the CbNs. Excitatory and inhibitory connections are depicted as round- and T-ends, respectively. Dashed lines represent connections not included in the model.

perimentally accessible; its architecture is well characterized; there is a clear feedback signal to the Purkinje cells (the cerebellar neurons principally involved in learning); its input-output function is relatively linear [36–38]; and molecular layer interneurons play a major role in shaping Purkinje cell responses [39–45, 47], and can influence climbing fiber-mediated dendritic calcium signals in Purkinje cells [46, 48, 61].

Both classic and more modern theoretical studies in the cerebellum have focused on the cerebellar cortex, modelling it as a one-layer feedforward network [62–66]. In this view, the parallel fibers project to Purkinje cells, and their synaptic weights are adjusted under the feedback signal from the climbing fibers. This picture, however, is an over-simplification, as Purkinje cells do not directly influence downstream structures. Instead, they project to the cerebellar nucleus, which constitutes the output of the cerebellum (see Fig. 5). The fact that Purkinje cells form a hidden layer, combined with the observed plasticity in the Purkinje cell to cerebellar nucleus synapses [67–71], means most learning rules tailored to one-layer networks, including the delta rule, cannot be used to train the network.

We propose instead that the cerebellum acts as a two layer DGN comprised of Purkinje cells as the first, hidden layer and the cerebellar nucleus as the second, output

layer (Fig. 5). Parallel fibers provide the input to both the input layer (Purkinje cells) and the gates, represented by molecular layer interneurons, that control learning in individual Purkinje cell dendrites. For the output layer of the DGN (which consists of one neuron), we use a non-gated rather than a gated neuron, as the unique biophysical features of cerebellar nuclear neurons allow them to integrate inputs approximately linearly [72]. The climbing fibers provide the feedback signal to Purkinje cells and cerebellar nuclear neurons. In our formulation, climbing fiber feedback signals the target, allowing each neuron to compute its own local error by comparing the target to its output ($r_{k,i}$). This formulation is a departure from the strict error-coding role that is traditionally attributed to climbing fibers, but is consistent with a growing body of evidence that climbing fibers signal a variety of sensorimotor and cognitive predictions [73].

DGNs can learn inverse kinematics

The cerebellum is thought to implement inverse motor control [74, 75]. We therefore applied our proposed DGN network to the SARCOS benchmark [76], which is an inverse kinematics dataset from a 7 degree-of-freedom robot arm (Fig. 6). The goal is to learn an inverse model, and predict 7 torques given the joint positions, velocities, and accelerations for each of the 7 joints (corresponding to a 21 dimensional input).

The target output, r^* , is the desired torque, given the 21-dimensional input. There are seven joints, so we train seven different networks, each with its own target output. We use DGN networks with 20 Purkinje cells, and minimize a quadratic loss (5). Since this is a relatively hard task, performance depends strongly on the number of branches. In Fig. 6 we plot the target torques for each joint (dots) along with the predictions of the DGN (lines; chosen for ease of comparison as there is no data between the points) for 500 branches. The lines follow the points reasonably closely, even when there are large fluctuations, indicating that the DGN is faithfully predicting torques. The performance of our network (mean squared error on test data in the original torque units) is comparable to that of most existing machine learning algorithms (Supplementary Table S1) while using fewer samples to learn. In Supplementary Fig. S2 we show the equivalent plot for 5, 50 and 5000 branches. Even at 5 branches performance is reasonable, while at 5000 we exceed the performance of almost all existing machine learning algorithms.

Vestibulo-ocular reflex, and adaptation to gain changes

To maintain a stable image on the retina during head movements, when an animal moves its head it moves its eyes in the opposite direction. This is known as the vestibulo-ocular reflex (VOR), and a key feature of it is that it's plastic: animals can adapt quickly when the relationship between the head movement and visual feedback is changed, as occurs as animals grow or are given corrective lenses. VOR gain adaptation relies critically on the cerebellum, and has been used to study cerebellar motor learning for decades [77–81]. This is an easy task to learn – almost any network, including a linear one, can achieve high performance on it. We consider it primarily because it is a very common cerebellar task.

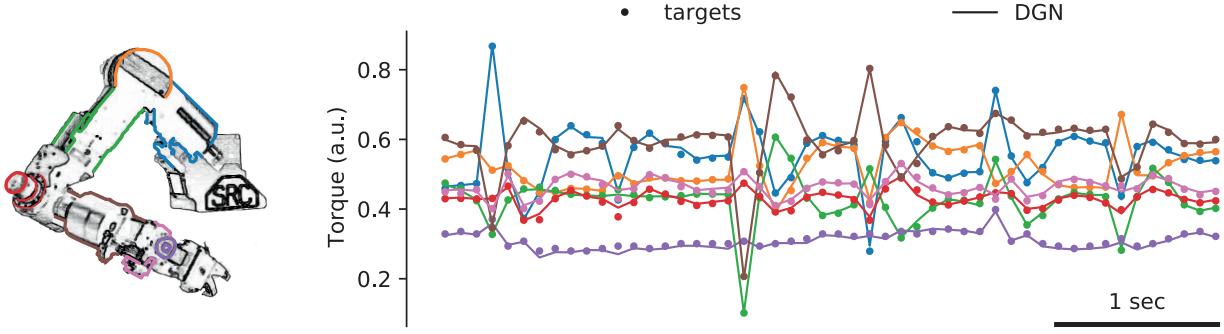


Figure 6: Sarcos experiment. DGNs can solve a challenging motor control task: predicting torques from the proprioceptive input. The data comes from a SARCOS dexterous robotic arm [76], pictured on the left. The inputs are position, velocity and acceleration of the 7 joints (a 21 dimensional variable); the targets are the desired torques (7 dimensional). Example targets (normalized to keep the training data between 0 and 1) are shown with dots, the lines are the output of our network. Performance is very good; only rarely is there a visible difference between the dots and the lines.

264 We applied our DGN network to the VOR in a regime where the gain occasionally
 265 changes abruptly. The gain, denoted G , is the ratio of the desired eye velocity to the
 266 head velocity (multiplied by -1 because the eyes and head move in opposite direction,
 267 to keep with the convention that the gain is reported as a positive number). When the
 268 gain is (artificially) changed, at first animals move their eyes at the wrong speed, but
 269 after about 15 minutes they learn to compensate [79, 80].

We trained our network on a head velocity signal of the form

$$s(t) = \sin(\omega_1 t) + \sin(\omega_2 t), \quad (7)$$

with $\omega_1 = 13.333$ and $\omega_2 = 20.733$ (corresponding to 2.12 and 3.30 Hz, respectively). This signal was chosen to mimic, approximately, the irregular head velocities encountered in natural viewing conditions. Following Clopath et. al. [82], we assumed that the Purkinje cells receive delayed versions of this signal. The i^{th} input signal, $x_i(t)$, which arrives via the parallel fibers, is modelled as

$$x_i(t) = s(t - \tau_i), \quad (8)$$

270 with delays, τ_i , spanning the range 50-300 ms. The cerebellum needs to compute the
 271 scaled version of the eye velocity: $r^*(t) = Gs(t)$ (as mentioned above, the actual eye
 272 movement is $-r^*(t)$, but we follow the standard convention). Learning was online, and
 273 we updated the weights every 500 ms, to approximately match the climbing fiber firing
 274 rate [83].

275 The DGN contained 20 Purkinje cells, with 10 branches each; these project to one
 276 output neuron (corresponding to the cerebellar nucleus), which was linear and not
 277 gated. As a baseline, we trained an MLP with the same number of weights (resulting

in 200 hidden neurons). We used quadratic loss for both the DGN and the MLP and, as in [82], we assumed $n = 100$ parallel fibers and a single output. Each branch received input from all 100 parallel fibers. Gating (Eq. (3)) was controlled by $x_i(t)$ (given in Eq. (8)), reflecting the parallel fiber influence on molecular layer interneurons (Fig. 5); see Methods for details. Given the timescale of the signal (2-3 Hz), any individual branch was gated on for about 500 ms at a time. The networks were pre-trained on a gain, G , of 1. We implemented four jump changes: first to 0.7, then back to 1.0, then to 1.3, and, finally, back to 1.0; in all cases, for 30 minutes at each gain (Fig. 7A).

Performance for both the DGN and the MLP were comparable and, after suitably adjusting the learning rates, the networks were able to learn in 15-20 minutes (Fig. 7A, B), consistent with learning times in behavioral experiments [79, 80]. Figure 7C shows the target and predicted head velocities immediately before and after each gain change. Not surprisingly, immediately after a gain change, the network produces output with the old gain.

Figure 7D shows the connection strengths between parallel fibers ($x_i(t)$, Eq. (8)) and Purkinje cells, after learning, as a function of the delay, τ_i . There are two notable features to these plots. First, the connectivity patterns are smooth. Second, although the DGN and the MLP solve the task equally well, there is a clear difference: for the MLP the connectivity patterns are highly stereotyped, while for the DGN they are far less so.

The smooth connectivity patterns, which are seen in both MLPs and DGNs, arise primarily because weights mediating inputs with similar delays have similar updates during learning. But there is another, somewhat technical, reason: the weights were initialized to small values. That's important because for most directions in weight space, changes in the weights have no effect on the loss. Component of the weights that lie in these "null" directions will, therefore, not change with learning. Small initial weights ensure that the components in the null directions start small, and the lack of learning in these directions means they stay small.

The difference in the connectivity patterns – stereotyped versus diverse – are due to the fact that MLPs are not gated whereas DGNs are. The smooth, stereotyped connectivity patterns seen in MLPs arise because all neurons receive similar input statistics, and so they find similar solutions. The more diverse connectivity patterns seen in DGNs arise because inputs to different branches are gated differently, and so different branches do not see the same input statistics.

What happens when the initial weights are large and initially random? In that case, because the weights don't change in the null directions, the final connectivity patterns are also not smooth – they're almost as noisy as the initial weights. Here as well, though, there are differences between DGNs and MLPs: for DGNs the noise rides on top of diverse connectivity patterns very similar to those in the top panels of Fig. 7D, while for MLPs the noise is unmodulated, and more or less white (see Supplementary Fig. S4).

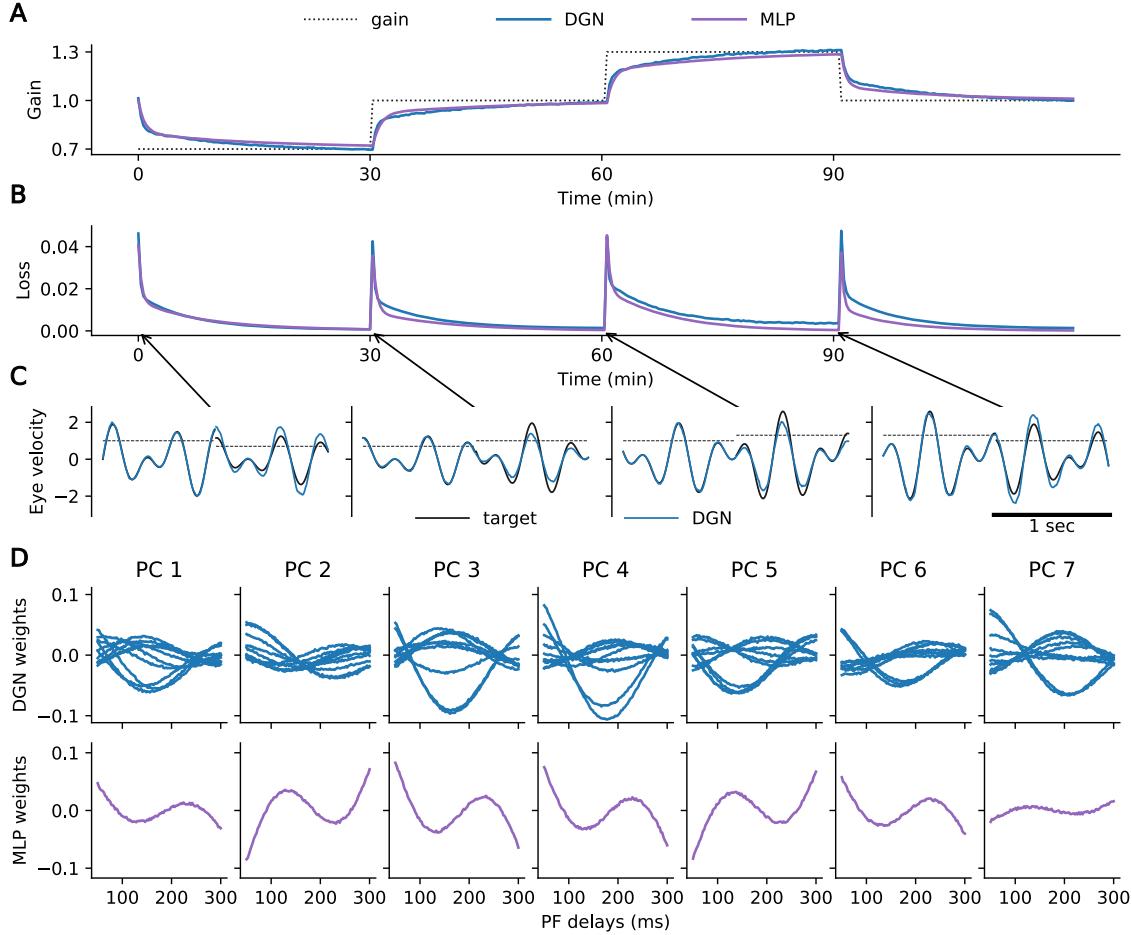


Figure 7: VOR adaptation task. We trained the networks on gain $G = 1$, then changed the gain every 30 minutes. Results are shown for the Dendritic Gated Network (DGN) and a multi-layer perceptron (MLP). **A.** Dashed lines are true gain versus time; blue and purple lines are gains computed by the DGN and MLP, respectively. For both networks, gains were inferred almost perfectly after 15-20 minutes. **B.** Performance, measured as mean squared error between the true angular velocity, $G_s(t)$ (Eq. (7)), and the angular velocity inferred by the networks. Same color code as panel A. **C.** Comparison of target angular velocity versus time (black) to that predicted by the DGN (blue). (A plot for the MLP is similar.) Before the gain change, the two are almost identical; immediately after the gain change, the network uses the previous gain. **D.** Top panel: Parallel fiber weights for the DGN network versus delay, τ_i (Eq. (8)). Each panel shows 10 branches; 5 Purkinje cells are shown (chosen randomly out of 20). The weights vary smoothly with delay. Bottom panel: MLP weight profile, except that dendritic branches are replaced by the whole neuron (all 100 parallel fibers). The weights again vary smoothly with delay, but their shapes are now highly stereotyped.

319 2.3 Testing predictions of the DGN in behaving animals.

320 A key feature of DGNs is that the gates should exert a local effect on the dendritic
321 branches of the principal neurons. When mapped onto the cerebellum, this suggests
322 that the molecular layer interneurons (MLIs) should inhibit the individual dendritic
323 branches of the Purkinje cells. We therefore tested whether individual Purkinje cell
324 dendrites can be inhibited by activity in MLIs. Previous *in vitro* work has demonstrated
325 that synaptic inhibition can locally inhibit calcium signals in Purkinje cell dendrites
326 [46], and *in vivo* work has shown that MLI activity can influence the variability of these
327 signals [48, 61], but it has not yet been shown whether such effects can be localized
328 to individual dendrites of Purkinje cells, and if so, what the spatial relationship of
329 this effect is with presynaptic interneurons. Using multi-plane 2-photon imaging in
330 awake PV-cre mice injected with the GCaMP7f virus, we could reliably record calcium
331 signals from individual MLI somata and the climbing fiber-evoked calcium signals in
332 the dendritic tree of Purkinje cells (Fig. 8A-B). With this approach, we were able
333 to identify a substantial proportion of interneurons (72/142 MLIs in 3 mice; 51%)
334 whose activation was associated with a significant decrease in climbing-fiber driven
335 calcium signals in the dendrites of at least one nearby Purkinje cell (Fig. S5A). Given
336 the axonal spread of MLIs [84, 85], the analysis was confined to single interneurons that
337 were located within 150 μm (rostrocaudally) and 50 μm (mediolaterally) from a given
338 Purkinje cell dendrite. The extent of suppression varied between PC dendrites and also
339 within individual PC dendrites recorded at different depths (Fig. 8C). In modulated
340 Purkinje cell dendrites, the degree of suppression was $17.4 \pm 0.5\%$ ($n = 133$ Purkinje
341 cells in 3 mice [range 6.6 to 53.5%]).

342 After identifying Purkinje cells whose global dendritic signals were inhibited when
343 nearby MLIs were active, we investigated the spatial extent of this inhibition within
344 Purkinje cell dendritic segments. We generated climbing fiber-evoked calcium signal
345 maps in Purkinje cell dendrites, both when MLIs were active and when they were
346 inactive (Fig. 8D, left). The difference revealed that MLI activation was associated
347 with local suppression of calcium signals in subregions of the dendrites (blue region in
348 Fig. 8E, left; Fig. S5B-E). To quantify the spatial extent of the suppression, we sub-
349 divided Purkinje cell dendritic regions into 1 μm segments to generate spatial activity
350 profiles in MLI-active and MLI-inactive conditions (Figs. 8D, right). Subtracting these
351 yielded a spatial difference trace (Figs. 8E, right). Aligning these segments across all
352 modulated PC dendrites allowed us to determine the average spatial profile of sup-
353 pressed dendritic segments (Figs. 8F). To identify false positives, for each Purkinje
354 cell dendrite analyzed we generated a shuffled difference trace, where MLI-active and
355 MLI-inactive traces were replaced with an equal mixture of the two conditions (odd-
356 even event split). We then identified suppressed regions in these shuffled difference
357 traces, and used the 95th percentile (4.6 μm) as the minimum length for a segment
358 to be considered significantly modulated. Across all experiments, we identified $n = 77$
359 significantly modulated segments ($N = 3$ mice) whose calcium signals were suppressed
360 by $42 \pm 2\%$ (mean \pm S.E.M.) (Fig. 8G). The spatial extent of MLI-gated inhibition
361 in these dendritic segments was $31 \pm 3 \mu\text{m}$ (mean \pm S.E.M.), accounting for $35 \pm 3\%$
362 (mean \pm S.E.M.) of the extent of the dendritic tree at the imaged plane (Fig. 8H-I).
363 These results demonstrate that MLIs can locally inhibit climbing fiber signals in the

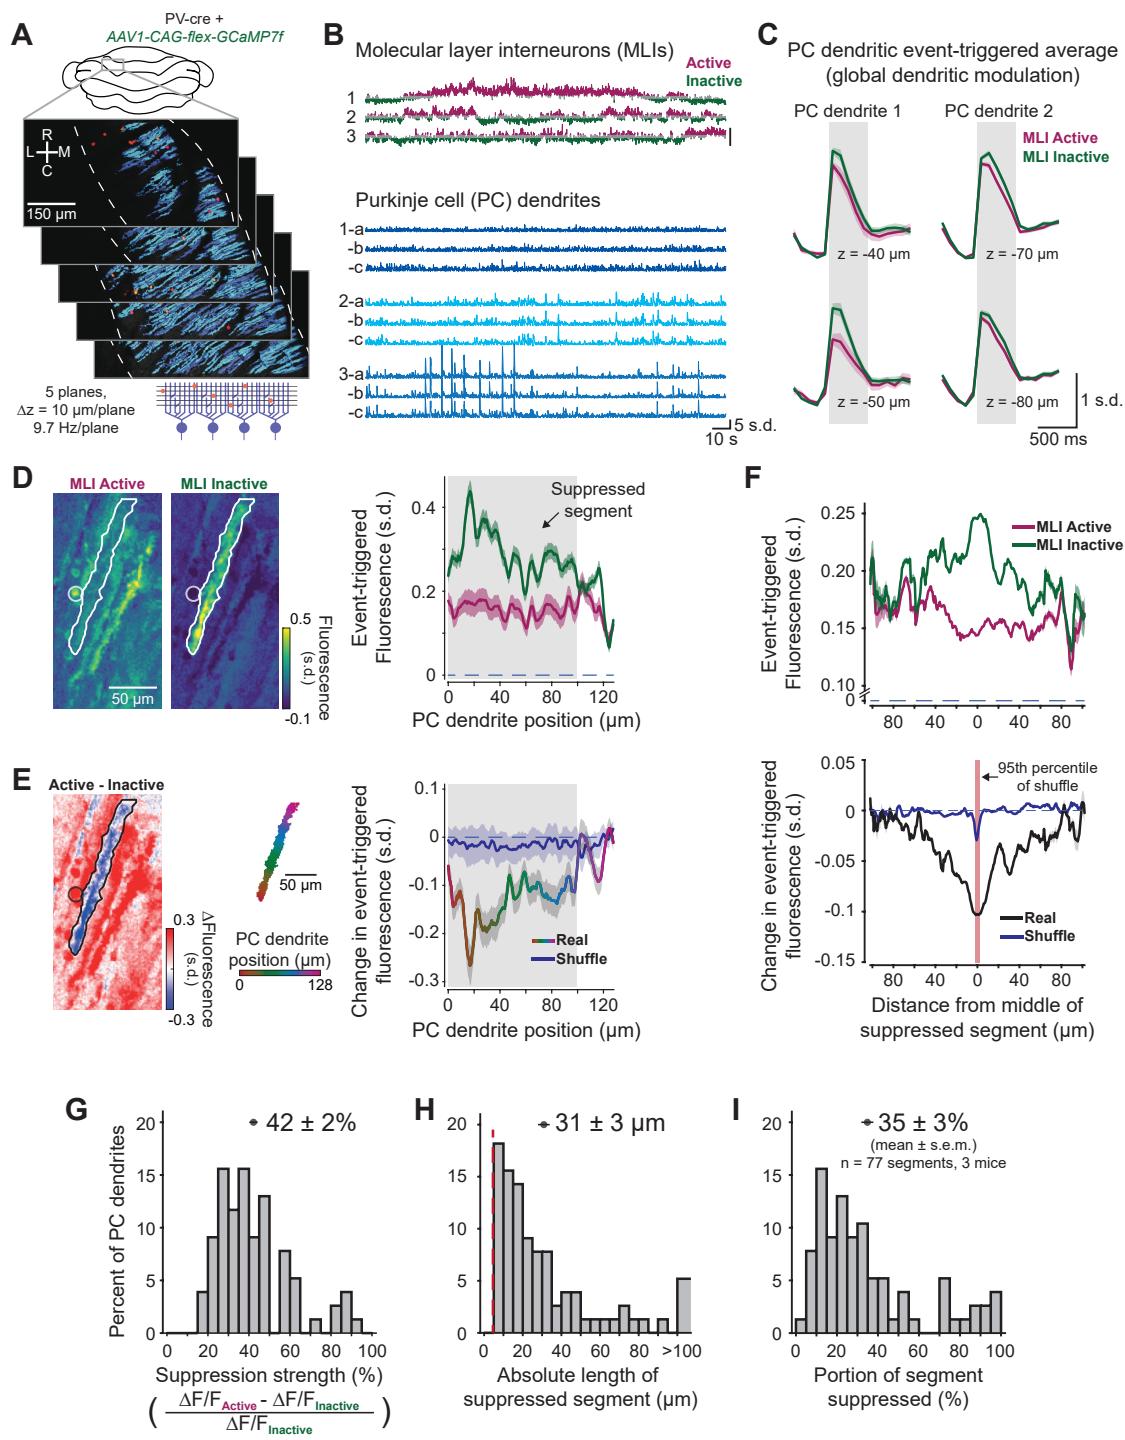


Figure 8: Dendritic gating of Purkinje cells by molecular layer interneurons *in vivo*.

Figure 8: **A.** Multi-plane 2-photon calcium imaging of molecular layer interneurons (red) and Purkinje cell dendrites (blue). **B.** Example traces of MLIs (top, active and inactive states in purple and green) and Purkinje cell dendrites recorded in multiple planes marked a, b, and c (bottom). **C.** Two examples of mean climbing fiber-evoked Ca^{2+} signal, in two planes, in MLI-active (purple) and inactive (green) states. **D.** Spatial event-triggered map of area surrounding Purkinje cell dendrites (contoured region of interest) when a nearby MLI (circle projected from different plane) was in an active state or inactive state (left). Spatial profile of event-triggered fluorescence is shown on the right. **E.** Difference heatmap image (left) and spatial profile (rainbow trace, right) of event-triggered fluorescence of same dendrite shown in panel D. Shuffled trace (blue) was computed on a split of odd and even event indices. **F.** Mean spatial profile of event-triggered fluorescence in MLI-active and MLI-inactive conditions (top) and difference traces (bottom) aligned to middle of suppressed segment (black trace, $n = 77$ regions from $N = 3$ recordings in 3 mice). Shuffled trace (blue) was computed, as above, on an odd-even split. Suppressed regions smaller than the 95th percentile of the shuffle data (red region, $4.6 \mu\text{m}$) were excluded. **G.** Histogram of strength of suppression calculated only within modulated segments. **H.** Histogram of suppressed segment lengths (red line shows 95th percentile of the shuffle). **I.** Histogram of suppressed segment lengths expressed as a percent of total segment length. Data are shown as mean \pm S.E.M.

364 Purkinje cell dendritic tree with a spatial extent that is comparable to the length of
365 individual spiny branchlets. This provides strong experimental evidence for a specific
366 biological implementation of one of the principal design features of the DGN, namely
367 local dendritic suppression of principal units by interneurons.

368 3 Discussion

369 Identifying the biologically plausible learning rules that mediate the modification of
370 connections in neural networks is a key goal of both experimental and theoretical
371 neuroscience. Here, we describe a new class of learning rules called Dendritic Gated
372 Networks (DGNs). Each unit in each layer of a DGN consists of dendritic branches
373 that are gated on and off by interneurons, and all units in all layers receive the same
374 feedback learning signal. We show that the DGN has key advantages over existing
375 learning algorithms, particularly in terms of learning speed and resilience to forgetting,
376 tested across a range of learning tasks. We also show using *in vivo* experiments that key
377 elements of the DGN architecture may be implemented in biological networks. These
378 results suggest that DGNs may be widely useful in the machine learning community,
379 and also suggest that this learning rule may be implemented in biological networks
380 such as the cerebellum and other neural circuits with a similar network architecture.

381 3.1 Comparison of DGNs to other learning algorithms

382 The DGN network architecture differs from traditional learning algorithms (e.g. back-
383 prop) as well as the algorithm on which it was based (Gated Linear Networks) in
384 several important ways. Traditional learning algorithms like backprop map input to
385 output in stages, with the input gradually transformed, until eventually, in the output
386 layer, the relevant features are easy to extract. There is certainly some evidence for
387 this hierarchical strategy being implemented in the brain. It is, for example, much
388 harder to extract which face a person is looking at from activity in visual area V1 than
389 in fusiform face area [86, 87]. While this strategy for computing is reasonable, it has
390 a downside: the relationship between activity in intermediate layers and activity in
391 the output layer is highly nontrivial, which makes it especially hard for the brain to
392 determine how weights in intermediate layers should change.

393 Despite the inherent complexity of this strategy, biologically plausible learning rules
394 that implement it have been proposed [15–27]. The DGN algorithm takes a different
395 approach from any of these. With this architecture, dendritic branches are gated on
396 and off via a random (and fixed) linear transformation of their input. The summed
397 activity of these branches forms the prediction of the neuron, which gets adjusted over
398 time via a delta rule. Consequently, all neurons predict the same target; and each layer
399 improves upon the predictions of the previous layer.

400 The DGN also differs from and improves over the algorithm on which it was based
401 – the Gated Linear Network (GLN) [29, 30]. In particular, the GLN requires a bank
402 of weights for each neuron, with the input choosing which one the neuron should use
403 – something that seems extremely difficult for the brain to implement. The DGN,
404 however, replaces the library of weights with gated dendritic branches, an innovation
405 essential for biological plausibility. Thus, although the DGN is conceptually related to
406 the GLN, from the point of view of neuroscience it has a critical new component which
407 makes it, unlike the GLN, relevant to the brain.

408 3.2 Implementation in cerebellar circuits

409 The architecture and exceptional efficiency of learning exhibited by DGNs suggests that
410 this algorithm may also be implemented in biological networks. Specifically, several key
411 features of the DGN are recapitulated in the functional architecture of the cerebellum.
412 First, the cerebellum receives a clear and global feedback signal in the form of the
413 climbing fiber input to Purkinje cells and cerebellar nuclear neurons that is the principal
414 driver of learning in the cerebellar circuit [88]. Second, the principal neurons of the
415 cerebellar cortex, Purkinje cells, exhibit linear encoding of their inputs due to their
416 high baseline firing rates and unique biophysical properties [36, 89]. Finally, molecular
417 layer interneurons, which are known to target dendritic branches of Purkinje cells, are
418 likely candidates to mediate branch-specific dendritic inhibition [46, 48, 90].

419 A key prediction of the DGN that would bolster its biological plausibility is that
420 interneurons should gate activity in single dendritic branches of principal cells. Here,
421 we provide the first *in vivo* evidence that molecular layer interneurons can produce
422 inhibition of dendritic calcium signals on the level of single dendritic branches in Purk-
423 inje cells, a longstanding, but until now untested, prediction of anatomical [91–94]

and theoretical [63, 90, 95] work. By simultaneously imaging dendritic calcium signals in Purkinje cells and activity in neighboring MLIs, we show that MLI activity can substantially decrease dendritic calcium signals in Purkinje cells. Previous *in vitro* work has shown that even modest inhibition of dendritic calcium signals (on order of 20%) can completely abolish cerebellar plasticity [96], suggesting that the suppression we observe is capable of abolishing learning. These MLI-driven effects were not distributed equally across PC dendritic arbors. Indeed, suppression of dendritic calcium signals was often restricted to individual dendritic branches of Purkinje cells, even as neighboring regions of the same dendritic arbor were unaffected. The profound local suppression of these climbing fiber-driven signals suggest that MLI-driven inhibition is also capable of suppressing parallel fiber-driven input to Purkinje cells, which are relatively much weaker [97]. Thus, it is likely that MLIs can gate both input and learning to single Purkinje cell dendritic branches. In summary, our demonstration that MLIs can modulate the Purkinje cell dendritic calcium signals on the spatial scale of a single dendritic spiny branchlet strongly supports the DGN gating prediction. We note that the binary on/off gating exhibited by the DGN is challenging to implement biologically; far more likely are softer gates, where the amount of gating is a function of the input. To determine how soft gates affected DGN performance, we performed simulations on the permuted MNIST and inverse kinematics tasks, and the results were virtually identical to the ones with on/off gates (data not shown), emphasizing the flexibility of DGN implementation in the brain. Our experimental data (together with previous work linking cerebellar functional architecture to features of the DGN) provides strong support for the idea that a DGN-like algorithm is implemented in cerebellar circuits.

The architecture of the DGN makes several further predictions about how the DGN may map onto the cerebellum. A key prediction that awaits experimental validation is that the activity of MLIs should depend predominantly on parallel fiber input and the input-output relationship of these interneurons should change very slowly relative to the timescale over which Purkinje cells learn, which can be measured in single trials [98]. Assessing the stability of the parallel fiber-mediated input-output relationship of MLIs is difficult because granule cells are known to exhibit learning-related changes in activation [99, 100]. Thus, further experiments to determine the stability of parallel fiber-MLI and parallel fiber-Purkinje simple spike input-output relationships over learning will need to account for changes in the firing patterns of cerebellar afferents. Another prediction of the DGN is that the parallel fiber connectivity pattern in the VOR or similar tasks should carry information about architecture. If parallel fiber-Purkinje cell connectivity is smooth and diverse (Fig. 7D, top panels), this supports a DGN-like implementation. If, on the other hand, connectivity is smooth and stereotyped (Fig. 7D, bottom panels) or noisy with no other temporal structure (Fig. S4), this argues for alternative implementations, such as an MLP. Distinguishing between these alternatives experimentally would require recording from populations of individual granule cells during VOR learning, then mapping the synaptic strength between those same granule cells and Purkinje cells.

466 3.3 DGNs in other neural circuits

467 While DGNs exhibit features that make them particularly well-suited for implementation
468 in the cerebellum, the general principles of this learning rule may be applicable
469 to a variety of brain circuits. In particular, the gating of dendritic signals that we
470 demonstrate in Purkinje cells may also be a feature of cortical networks, given the
471 branch-specific innervation of interneuron axons that has been documented in many
472 cortical circuits [101]. Such generalization would require some modifications to imple-
473 ment the algorithm; for instance, the learning rule will have to change because activity
474 in the cortex is far from linear.

475 3.4 Conclusions

476 In summary, Dendritic Gated Networks are strong candidates for biological networks
477 – and not just in the cerebellum; they could be used anywhere there is approximately
478 feedforward structure. They come with two highly desirable features: rapid, data-
479 efficient learning, and biologically plausible learning. Furthermore, they suggest a
480 novel role for inhibitory neurons, which is that they are used for gating dendritic
481 branches. We anticipate that the strong, experimentally testable, predictions may
482 inspire investigations in many brain circuits where rapid learning may invoke a DGN
483 algorithm.

484 4 Methods

485 4.1 Classification tasks

486 The network we use in our model is described in Eqs. (1) and (3), and the learning rules
487 are given in Eq. (4). For regression (VOR and inverse kinematics), we use the identity
488 function for ϕ (Eq. (1)), and the square loss (Eq. (5)), resulting in the update rule
489 given in Eq. (6). For classification (the simple task described in Fig. 3 and permuted
490 MNIST), the network computes probabilities, so ϕ needs to be bounded, and a square
491 loss is not appropriate. Here we provide details for this case.

492 For classification we use a standard sigmoid function, $\sigma(z) = e^z/(1 + e^z)$, albeit
493 modified slightly,

$$\phi(z) = \text{CLIP}_\epsilon^{1-\epsilon}(\sigma(z)) \quad (9)$$

494 where $\text{CLIP}_a^b(\cdot)$ clips values between a and b (so the right hand side is zero if $\sigma(z)$ is
495 smaller than ϵ or larger than $1 - \epsilon$). Clipping is used for bounding the loss as well as
496 the gradients, which helps with numerical stability. It also enables a worst-case regret
497 analysis [29, 30]. We set ϵ to 0.01, so neural activity lies between 0.01 and 0.99.

498 The square loss is not appropriate in this case, so instead we use the binary cross-
499 entropy loss: the loss of neuron i in layer k is given by

$$\ell(r^*, r_{k,i}) = -r^* \log r_{k,i} - (1 - r^*) \log (1 - r_{k,i}). \quad (10)$$

498 Consequently, the update rule for the weights, Eq. (4), is (after a small amount of
 499 algebra)

$$\Delta w_{k,ij}^b = \eta g_{k,i}^b(\mathbf{x}) \mathbb{1}(\epsilon < r_{k,i} < 1 - \epsilon) (r^* - r_{k,i}) h_{k-1,j} \quad (11)$$

500 where $\mathbb{1}(\cdot)$ is 1 when its argument is true and 0 otherwise. The fact that learning is
 501 zero when $r_{k,i}$ is outside the range $[\epsilon, 1 - \epsilon]$ follows because $d\phi(z)/dz = 0$ when z is
 502 outside this range (see Eq. (9)). This ensures that learning saturates when weights
 503 become too large (either positive or negative). However, this can cause problems if the
 504 output is very wrong; that is, when $r^* = 1$ and $r_{k,i} < \epsilon$ or $r^* = 0$ and $r_{k,i} > 1 - \epsilon$. To
 505 address this, we allow learning in this regime. We can do that compactly by changing
 506 the learning rule to

$$\Delta w_{k,ij}^b = \eta g_{k,i}^b(\mathbf{x}) \mathbb{1}(|r^* - r_{k,i}| > \epsilon) (r^* - r_{k,i}) h_{k-1,j}. \quad (12)$$

507 Essentially, this rule says: stop learning when $r_{k,i}$ is within ϵ of r^* .

508 For a compact summary of the equations (given as pseudocode), see Supplementary
 509 Algorithms S1 and S2.

510 4.2 Simulations

511 Here we provide details of our simulations. Simulations were written using JAX [102],
 512 the DeepMind JAX Ecosystem [103], and Colab [104].

513 **Permuted MNIST.** We adopt the pixel-permuted MNIST benchmark [57, 60],
 514 which is a sequence of MNIST digit classification tasks with different pixel permutations.
 515 Each task consists of 60,000 training images and 10,000 test images; all images
 516 are deskewed. Models are trained sequentially across 10 tasks, performing a single pass
 517 over all 60,000 training examples for each of the tasks. We provide the implementation
 518 details below; the parameters swept during a grid search are given in Supplementary
 519 Table S2.

520 *DGN.* We use networks composed of 100 and 20 units in the hidden layers and a
 521 single linear neuron for the output. All neurons in the hidden layers have 10 dendritic
 522 branches. The output of the network is determined by the last neuron. MNIST has 10
 523 classes, each corresponding to a digit. Therefore, we use 10 DGN networks, each en-
 524 coding the probability of a distinct class. These networks are updated during training
 525 using a learning rate $\eta = 10^{-2}$. During testing, the class with the maximum proba-
 526 bility is chosen. Images are scaled and shifted so that the input range is $[-1, 1]$. The
 527 gating vectors, $\mathbf{v}_{k,i}^b$, are chosen randomly on the unit sphere, which can be achieved
 528 by sampling from an isotropic Normal distribution and then dividing by the L2 norm.
 529 The biases, $\theta_{k,i}^b$, are drawn independently from a zero mean Gaussian with standard
 530 deviation 0.05.

531 *MLP and EWC.* We use a ReLu network with 1000 and 200 neurons in the hidden
 532 layers and 10 linear output units with cross entropy loss. In this setting, the MLP and

533 EWC have the same number of neurons as the DGN, but fewer connections. We use the
534 ADAM optimization method [105] with a learning rate $\eta = 10^{-4}$ (see Supplementary
535 Table S2 for details of the hyperparameter optimization), in conjunction with dropout.
536 We use mini-batches of 20 data points. For EWC, we draw 100 samples for computing
537 the Fisher matrix diagonals and set the regularization constant to 10^3 .

538 **Inverse Kinematics.** Each DGN network has 20 Purkinje cells and one linear, non-
539 gated, output neuron, and we vary the number of branches. We use a quadratic loss,
540 as in Eq. (5), a learning rate $\eta = 10^{-5}$, and we run for 2000 epochs (2000 passes over
541 the dataset). The inputs are centered at 0 and scaled to unit variance per dimension,
542 and the targets are scaled so that they lie between 0 and 1. The reported MSEs are
543 computed on the test set based on inverse transformed predictions (thus undoing the
544 target scaling). The gating parameters are chosen in the same way as for the MNIST
545 simulations described above.

546 We discovered that the training set of the SARCOS dataset (downloaded from
547 <http://www.gaussianprocess.org/gpml/data/> on 15 December 2020) includes test
548 instances. To the best of our knowledge, other recent studies using the SARCOS
549 dataset [106, 107] reported results with this

550 train/test setting. This means that the reported errors are measures of capacity
551 rather than generalization. We compare the performance of DGN against the best
552 known SARCOS results in Supplementary Table S1 using the existing train/test split.

553 **VOR.** The gating parameters $v_{k,ij}^b$ and $\theta_{k,i}^b$ (Eq. (3)), were drawn independently
554 from the standard normal distribution. The learning rates were $\eta = 10^{-5}$ for the DGN
555 and $\eta = 0.03$ for the MLP.

556 4.3 Animal experiments

557 **Animal housing and surgery** All animal procedures were approved by the local
558 Animal Welfare and Ethical Review Board and performed under license from the
559 UK Home Office in accordance with the Animals (Scientific Procedures) Act 1986
560 and generally followed procedures described previously [108]. Briefly, we used PV-Cre
561 mice (B6;129P2-Pvalbtm1(cre)Arbr/J) [109] maintained on a C57/BL6 background.
562 Mice were group housed before and after surgery and maintained on a 12:12 day-night
563 cycle. Surgical procedures were similar to those described in [108], except that we in-
564 jected Cre-dependent GCaMP7f (pGP-AAV-CAG-FLEX-jGCaMP7f-WPRE [serotype
565 1]; [110]) diluted from its stock titer to a final concentration of 3×10^{11} GC/ml ($\sim 1:25$).
566 After mice had recovered from surgery, they were acclimated to the recording setup
567 and expression-checked before beginning recordings.

568 **Two-photon calcium imaging data acquisition and processing** Imaging
569 experiments were performed using a 16x/0.8 NA objective (Nikon) mounted on a Sutter
570 MOM microscope equipped with the Resonant Scan Box module. A Ti:Sapphire laser
571 tuned to 930 nm (Mai Tai, Spectra Physics) was raster scanned using a resonant
572 scanning galvanometer (8 kHz, Cambridge Technologies) and images were collected at

573 512x256 pixel resolution over fields of view of 670x335 μm per plane at an average
574 power of 30-70 mW. Volumetric imaging across 5 planes spaced by 10 μm (40 μm
575 depth range per recording, total depth ranging 10-70 μm below pial surface) were
576 performed using a P-726 PIFOC High-Load Objective Scanner (Physik Instruments)
577 at an effective volume rate of 9.7 Hz. The microscope was controlled using ScanImage
578 (Version 5.6, Vidrio Technologies) and tilted to 12.5 degrees such that the objective
579 was orthogonal to the surface of the brain and coverglass.

580 Regions of interest (ROIs) corresponding to single MLI somata and Purkinje cell
581 dendrites, which were easily distinguishable based on their shape, were identified using
582 Suite2p software [111] for initial source extraction and custom-written software for
583 subsequent analyses. Fluorescence time series were computed as $(F - F_0)/F_0$ where
584 F was the signal measured at each point in time and F_0 is the 8th percentile of a
585 rolling average baseline surrounding each data point (2000 frames for MLIs and 10
586 frames for Purkinje cell dendrites). A neuropil correction coefficient of 0.5 (50 percent
587 of neuropil signal output from Suite2p) was applied to MLI ROIs. A range of baseline
588 durations and neuropil correction coefficients were tested and varying these parameters
589 did not alter the main findings. Following these calculations, fluorescence signals were
590 z-scored to facilitate comparisons across neurons. Signals from the same Purkinje cell
591 dendrites recorded in multiple planes were identified based on a correlation threshold
592 (>0.5 , followed by manual confirmation) and analyzed independently for dendritic
593 modulation experiments. Event times in dendrites were detected using MLspike [112].

594 **MLI gating of Purkinje cell dendritic signals** For dendritic modulation ex-
595 periments, active and inactive MLI states were defined as imaging frames where activity
596 in an MLI deviated more than 0.5 standard deviations above or below the mean, re-
597 spectively (Fig. S5). Using a higher threshold yielded similar results but resulted in
598 fewer dendritic events in each condition. After identifying these states, we compared
599 the magnitudes of isolated dendritic events in these two conditions for Purkinje cells
600 within a $300 \times 100 \mu\text{m}$ ellipse centered on each MLI whose major axis is parallel to that
601 of Purkinje cell dendrites in the field of view (approximately rostrocaudal). Ellipse di-
602 mensions were chosen to approximate the known rostro-caudal and mediolateral spread
603 of MLI axons [84,85]. Only isolated events in Purkinje cell dendrites (defined as those
604 that occurred more than 500 ms before and after any other events) were analyzed, and
605 fluorescence event magnitudes were calculated over the 5 frames (~ 500 ms) after each
606 event for initial identification of MLI-modulated dendrites. Because analysis of each
607 recording involved many thousands of comparisons, we assessed significance differences
608 between Purkinje cell dendrite event magnitudes in MLI active and inactive states with
609 a significance threshold of 0.05 that was corrected for multiple comparisons using false
610 discovery rate threshold of 5% [113].

611 Motion-corrected fluorescence movies that were used for pixel-wise analysis of den-
612 dritic subregions were pre-processed by first correcting for slow fluctuations in fluores-
613 cence, which was done by computing $(F - F_0)/F_0$ where F was the signal measured at
614 each point in time and F_0 is the 8 percentile of a rolling average baseline surrounding
615 each data point (2000 frames), and then z -scoring. Purkinje cell dendritic ROIs defined
616 in Suite2p were segmented into 1 μm increments by fitting a 4th degree polynomial

617 to each ROI, grouping ROI pixels closest to regular spaced points along this fit line,
618 computing a weighted average of these pixels based on the pixel weights assigned by
619 Suite2p, and smoothing over $5\text{ }\mu\text{m}$. Activity profiles in MLI active and inactive states
620 were subtracted and used to generate spatial suppression profiles for each Purkinje
621 cell dendrite and error bars were generated from the summed variances in active and
622 inactive conditions. Shuffled distributions were generated by replacing active and in-
623 active conditions with an odd-even event split of each of these groups, yielding two
624 distributions comprising of 1/2 MLI-active and 1/2 MLI-inactive events for each den-
625 dritic segment. Significantly modulated dendritic segments were defined as the longest
626 region for each dendritic segment in which the 95% confidence interval of the difference
627 trace was less than zero. To account for false positive dendritic segments that would
628 be identified by finding minima in noise, we performed this identification procedure on
629 our shuffled data and excluded identified dendritic segments in our real data that were
630 smaller than the 95th percentile of these fictive segments ($4.6\text{ }\mu\text{m}$).

631 **Code availability.** We provide pseudo code in Supplementary Algorithms S1 and S2.
632 A simple python implementation can be accessed via https://github.com/deepmind/deepmind-research/blob/master/gated_linear_networks/colabs/dendritic_gated_network.ipynb.

635 **Data availability.** The data that support the findings of this study are available
636 from the corresponding authors upon reasonable request. Additional analysis made
637 use of standard publicly available benchmarks including MNIST [114] and SARCOS
638 (<http://www.gaussianprocess.org/gpml/data/>).

639 **Acknowledgements.** We thank Timothy Lillicrap, Gregory Wayne, Eszter Vértes,
640 and Brendan Bicknell for valuable feedback. Michael Häusser is supported by the
641 Wellcome Trust and the European Research Council. Peter Latham is supported by
642 the Gatsby Charitable Foundation and the Wellcome Trust.

643 **Author contributions.** ES and AGB developed the computational model with
644 advice from JV, CC, PEL, DB, and MHut. AGB, ES, and SK performed simulation
645 experiments and analysis with advice from CC and PEL. DK and MBea acquired and
646 analyzed neuronal data with advice from MHäu. PEL, AGB, ES, and DK wrote the
647 paper with help from all other authors. AGB and ES managed the project with support
648 from MBot, CC, JV, and DB.

649 **Competing Interests.** The authors declare no competing interests.

650 References

- 651 [1] Burke, C. J. *et al.* Layered reward signalling through octopamine and dopamine
652 in drosophila. *Nature* **492**, 433–437 (2012).

- 653 [2] Liu, C. *et al.* A subset of dopamine neurons signals reward for odour memory in
654 drosophila. *Nature* **488**, 512–516 (2012).
- 655 [3] Schmidhuber, J. Deep learning in neural networks: An overview. *Neural networks*
656 **61**, 85–117 (2015).
- 657 [4] Grossberg, S. Competitive learning: From interactive activation to adaptive
658 resonance. *Cogn. Sci.* **11**, 23–63 (1987).
- 659 [5] Crick, F. The recent excitement about neural networks. *Nature* **337**, 129–132
660 (1989).
- 661 [6] Roelfsema, P. R. & Holtmaat, A. Control of synaptic plasticity in deep cortical
662 networks. *Nature Reviews Neuroscience* **19**, 166 (2018).
- 663 [7] Whittington, J. C. & Bogacz, R. Theories of error back-propagation in the brain.
664 *Trends in cognitive sciences* **23**, 235–250 (2019).
- 665 [8] Lillicrap, T. P., Santoro, A., Marris, L., Akerman, C. J. & Hinton, G. Backprop-
666 agation and the brain. *Nature Reviews Neuroscience* **21**, 335–346 (2020).
- 667 [9] Kelley, H. J. Gradient theory of optimal flight paths. *Ars Journal* **30**, 947–954
668 (1960).
- 669 [10] Linnainmaa, S. The representation of the cumulative rounding error of an al-
670 gorithm as a taylor expansion of the local rounding errors. *Master's Thesis (in*
671 *Finnish)*, Univ. Helsinki 6–7 (1970).
- 672 [11] Rumelhart, D. E., , G. E. & Williams, R. J. Learning representations by back-
673 propagating errors. In Anderson, J. A. & Rosenfeld, E. (eds.) *Neurocomputing:*
674 *Foundations of Research*, 696–699 (MIT Press, Cambridge, MA, USA, 1988).
- 675 [12] French, R. Catastrophic forgetting in connectionist networks. *Trends in cognitive*
676 *sciences* **3**, 128–135 (1999).
- 677 [13] Herzfeld, D. J., Vaswani, P. A., Marko, M. K. & Shadmehr, R. A memory of
678 errors in sensorimotor learning. *Science* **345**, 1349–1353 (2014).
- 679 [14] Wolpert, D. M. & Flanagan, J. R. Computations underlying sensorimotor learn-
680 ing. *Current opinion in neurobiology* **37**, 7–11 (2016).
- 681 [15] Lillicrap, T. P., Cownden, D., Tweed, D. B. & Akerman, C. J. Random synap-
682 tic feedback weights support error backpropagation for deep learning. *Nature*
683 *Communications* **7**, 13276 (2016). URL <http://www.nature.com/articles/ncomms13276>.
- 685 [16] Akrout, M., Wilson, C., Humphreys, P., Lillicrap, T. & Tweed, D. B. Deep
686 learning without weight transport. In *Advances in Neural Information Processing*
687 *Systems*, vol. 6, e22901 (2017).
- 688 [17] Guerguiev, J., Lillicrap, T. P. & Richards, B. A. Towards deep learning with
689 segregated dendrites. *eLife* **6**, e22901 (2019).
- 690 [18] Sacramento, J., Costa, R. P., Bengio, Y. & Senn, W. Dendritic cortical mi-
691 crocircuits approximate the backpropagation algorithm. In *Advances in Neural*
692 *Information Processing Systems*, 8721–8732 (2018).

- 693 [19] Payeur, A., Guergiev, J., Zenke, F., Richards, B. A. & Naud, R. Burst-
694 dependent synaptic plasticity can coordinate learning in hierarchical circuits.
695 *bioRxiv* (2020). URL <https://www.biorxiv.org/content/early/2020/03/31/2020.03.30.015511>.
- 697 [20] Samadi, A., Lillicrap, T. P. & Tweed, D. B. Deep learning with dynamic spiking
698 neurons and fixed feedback weights. *Neural Computation* **29**, 578–602 (2017).
- 699 [21] Belilovsky, E., Eickenberg, M. & Oyallon, E. Greedy layerwise learning can
700 scale to ImageNet. In *Proceedings of the 36th International Conference on Ma-*
701 *chine Learning*, vol. 97, 583–593 (PMLR, 2019). URL <http://proceedings.mlr.press/v97/belilovsky19a.html>.
- 702 [22] Nøkland, A. Direct feedback alignment provides learning in deep neural networks.
703 In *Proceedings of the 30th International Conference on Neural Information Pro-*
704 *cessing Systems*, NIPS’16, 1045–1053 (Curran Associates Inc., Red Hook, NY,
705 USA, 2016).
- 706 [23] Nøkland, A. & Eidnes, L. H. Training Neural Networks with Local Error Sig-
707 nals. In *International Conference on Machine Learning*, 4839–4850 (PMLR,
708 2019). URL <http://proceedings.mlr.press/v97/nokland19a.html>. ISSN:
709 2640-3498.
- 710 [24] Löwe, S., O’Connor, P. & Veeling, B. Putting an end to end-to-end: Gradient-
711 isolated learning of representations. In *Advances in Neural Information Process-*
712 *ing Systems*, 3033–3045 (2019).
- 713 [25] Pogodin, R. & Latham, P. E. Kernelized information bottleneck leads to bio-
714 logically plausible 3-factor hebbian learning in deep networks. *arXiv preprint arXiv:2006.07123* (2020). 2006.07123.
- 715 [26] Podlaski, W. F. & Machens, C. K. Biological credit assignment through dynamic
716 inversion of feedforward networks (2020). 2007.05112.
- 717 [27] Golkar, S., Lipshutz, D., Bahroun, Y., Sengupta, A. M. & Chklovskii, D. B. A
718 biologically plausible neural network for local supervision in cortical microcircuits
719 (2020). 2011.15031.
- 720 [28] Clark, D. G., Abbott, L. & Chung, S. Credit assignment through broadcasting
721 a global error vector. *arXiv preprint arXiv:2106.04089* (2021).
- 722 [29] Veness, J. *et al.* Online learning with gated linear networks. *arXiv preprint arXiv:1712.01897* (2017).
- 723 [30] Veness, J. *et al.* Gated linear networks. *Proceedings of the AAAI Conference on Artificial Intelligence (To Appear)* (2021).
- 724 [31] Budden, D. *et al.* Gaussian gated linear networks. In *Advances in Neural Infor-*
725 *mation Processing Systems* (2020).
- 726 [32] Sezener, E., Hutter, M., Budden, D., Wang, J. & Veness, J. Online learning in
727 contextual bandits using gated linear networks. In *Advances in Neural Infor-*
728 *mation Processing Systems* (2020).

- 733 [33] Wang, J., Sezener, E., Budden, D., Mutter, M. & Veness, J. A combinatorial
734 perspective on transfer learning. In *Advances in Neural Information Processing*
735 *Systems* (2020).
- 736 [34] Widrow, B. & Hoff, M. E. Adaptive switching circuits. In *1960 IRE WESCON*
737 *Convention Record, Part 4*, 96–104 (IRE, New York, 1960).
- 738 [35] Tsuda, B., Tye, K. M., Siegelmann, H. T. & Sejnowski, T. A modeling framework
739 for adaptive lifelong learning with transfer and savings through gating in the
740 prefrontal cortex. *Proceedings of the National Academy of Sciences* **117**, 29872–
741 29882 (2020).
- 742 [36] Llinás, R. & Sugimori, M. Electrophysiological properties of in vitro purkinje cell
743 somata in mammalian cerebellar slices. *The Journal of physiology* **305**, 171–195
744 (1980).
- 745 [37] Walter, J. T. & Khodakhah, K. The linear computational algorithm of cerebellar
746 purkinje cells. *Journal of Neuroscience* **26**, 12861–12872 (2006).
- 747 [38] Chen, S., Augustine, G. J. & Chadderton, P. Serial processing of kinematic
748 signals by cerebellar circuitry during voluntary whisking. *Nature communications*
749 **8**, 1–13 (2017).
- 750 [39] Eccles, J. Functional meaning of the patterns of synaptic connections in the
751 cerebellum. *Perspectives in biology and medicine* **8**, 289–310 (1965).
- 752 [40] Dizon, M. J. & Khodakhah, K. The role of interneurons in shaping purkinje
753 cell responses in the cerebellar cortex. *Journal of Neuroscience* **31**, 10463–10473
754 (2011).
- 755 [41] Brown, A. M. *et al.* Molecular layer interneurons shape the spike activity of
756 cerebellar purkinje cells. *Scientific reports* **9**, 1–19 (2019).
- 757 [42] Andersen, P., Eccles, J. & Voorhoeve, P. Postsynaptic inhibition of cerebellar
758 purkinje cells. *Journal of neurophysiology* **27**, 1138–1153 (1964).
- 759 [43] Mittmann, W., Koch, U. & Häusser, M. Feed-forward inhibition shapes the
760 spike output of cerebellar purkinje cells. *The Journal of physiology* **563**, 369–378
761 (2005).
- 762 [44] Häusser, M. & Clark, B. A. Tonic synaptic inhibition modulates neuronal output
763 pattern and spatiotemporal synaptic integration. *Neuron* **19**, 665–678 (1997).
- 764 [45] Oldfield, C. S., Marty, A. & Stell, B. M. Interneurons of the cerebellar cortex
765 toggle purkinje cells between up and down states. *Proceedings of the National*
766 *Academy of Sciences* **107**, 13153–13158 (2010).
- 767 [46] Callaway, J. C., Lasser-Ross, N. & Ross, W. N. Ipsps strongly inhibit climbing
768 fiber-activated $[Ca^{2+}]_i$ increases in the dendrites of cerebellar purkinje neurons.
769 *Journal of Neuroscience* **15**, 2777–2787 (1995).
- 770 [47] Wulff, P. *et al.* Synaptic inhibition of purkinje cells mediates consolidation of
771 vestibulo-cerebellar motor learning. *Nature neuroscience* **12**, 1042–1049 (2009).
- 772 [48] Gaffield, M. A., Rowan, M. J., Amat, S. B., Hirai, H. & Christie, J. M. Inhibi-
773 tion gates supralinear Ca^{2+} signaling in purkinje cell dendrites during practiced
774 movements. *Elife* **7**, e36246 (2018).

- 775 [49] Carpenter, G. A. & Grossberg, S. The art of adaptive pattern recognition by a
776 self-organizing neural network. *Computer* **21**, 77–88 (1988).
- 777 [50] McCloskey, M. & Cohen, N. J. Catastrophic interference in connectionist net-
778 works: The sequential learning problem. *Psychology of Learning and Motiva-
779 tion* **24**, 109 – 165 (1989). URL <http://www.sciencedirect.com/science/article/pii/S0079742108605368>.
- 780 [51] Robins, A. V. Catastrophic forgetting, rehearsal and pseudorehearsal. *Connect.
781 Sci.* **7**, 123–146 (1995).
- 782 [52] Caruana, R. Multitask learning. *Machine Learning* **28**, 41–75 (1997). URL
783 <https://doi.org/10.1023/A:1007379606734>.
- 784 [53] Rebuffi, S.-A., Kolesnikov, A., Sperl, G. & Lampert, C. H. icarl: Incremental
785 classifier and representation learning. *2017 IEEE Conference on Computer Vi-
786 sion and Pattern Recognition (CVPR)* (2017). URL <http://dx.doi.org/10.1109/CVPR.2017.587>.
- 787 [54] Donahue, J. *et al.* Decaf: A deep convolutional activation feature for generic
788 visual recognition. *CoRR* **abs/1310.1531** (2013). URL <http://arxiv.org/abs/1310.1531>. 1310.1531.
- 789 [55] Razavian, A. S., Azizpour, H., Sullivan, J. & Carlsson, S. Cnn features off-
790 the-shelf: An astounding baseline for recognition. *2014 IEEE Conference on
791 Computer Vision and Pattern Recognition Workshops* (2014). URL <http://dx.doi.org/10.1109/CVPRW.2014.131>.
- 792 [56] Girshick, R., Donahue, J., Darrell, T. & Malik, J. Rich feature hierarchies for
793 accurate object detection and semantic segmentation. *2014 IEEE Conference on
794 Computer Vision and Pattern Recognition* (2014). URL <http://dx.doi.org/10.1109/CVPR.2014.81>.
- 795 [57] Kirkpatrick, J. *et al.* Overcoming catastrophic forgetting in neural net-
796 works. *Proceedings of the National Academy of Sciences* **114**, 3521–3526 (2017).
797 URL <https://www.pnas.org/content/114/13/3521>. <https://www.pnas.org/content/114/13/3521.full.pdf>.
- 798 [58] Zenke, F., Poole, B. & Ganguli, S. Continual learning through synaptic intelli-
799 gence. In *Proceedings of the 34th International Conference on Machine Learning - Volume 70*, ICML’17, 3987–3995 (JMLR.org, 2017).
- 800 [59] Schwarz, J. *et al.* Progress &; compress: A scalable framework for continual
801 learning. In Dy, J. & Krause, A. (eds.) *Proceedings of the 35th International
802 Conference on Machine Learning*, vol. 80 of *Proceedings of Machine Learning Re-
803 search*, 4528–4537 (PMLR, Stockholm, Sweden, 2018). URL
804 <http://proceedings.mlr.press/v80/schwarz18a.html>.
- 805 [60] Goodfellow, I. J., Mirza, M., Xiao, D., Courville, A. & Bengio, Y. An empirical
806 investigation of catastrophic forgetting in gradient-based neural networks (2013).
807 1312.6211.
- 808 [61] Kitamura, K. & Häusser, M. Dendritic calcium signaling triggered by sponta-
809 neous and sensory-evoked climbing fiber input to cerebellar purkinje cells in vivo.
810 *Journal of Neuroscience* **31**, 10847–10858 (2011).
- 811
- 812
- 813
- 814
- 815
- 816
- 817

- 818 [62] Marr, D. A theory of cerebellar cortex. *The Journal of Physiology* **202**, 437–470.1
819 (1969). URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1351491/>.
- 820 [63] Albus, J. S. A theory of cerebellar function. *Mathematical Biosciences* **10**,
821 25–61 (1971). URL <http://www.sciencedirect.com/science/article/pii/0025556471900514>.
- 822 [64] Ito, M. & Kano, M. Long-lasting depression of parallel fiber-purkinje cell trans-
823 mission induced by conjunctive stimulation of parallel fibers and climbing fibers
824 in the cerebellar cortex. *Neuroscience Letters* **33**, 253 – 258 (1982). URL
825 <http://www.sciencedirect.com/science/article/pii/0304394082903809>.
- 826 [65] Narain, D., Remington, E. D., Zeeuw, C. I. D. & Jazayeri, M. A cerebellar mech-
827 anism for learning prior distributions of time intervals. *Nature Communications*
828 **9**, 469 (2018).
- 829 [66] Shadmehr, R. Population coding in the cerebellum and its implications for learn-
830 ing from error. *bioRxiv* (2020).
- 831 [67] Pedroarena, C. M. & Schwarz, C. Efficacy and short-term plasticity at gabaergic
832 synapses between purkinje and cerebellar nuclei neurons. *Journal of Neurophys-
833 iology* **89**, 704–715 (2003).
- 834 [68] Pedroarena, C. M. Short-term plasticity at purkinje to deep cerebellar nuclear
835 neuron synapses supports a slow gain-control mechanism enabling scaled linear
836 encoding over second-long time windows. *bioRxiv* 749259 (2019).
- 837 [69] Morishita, W. & Sastry, B. Postsynaptic mechanisms underlying long-term de-
838 pression of gabaergic transmission in neurons of the deep cerebellar nuclei. *Jour-
839 nal of Neurophysiology* **76**, 59–68 (1996).
- 840 [70] Aizenman, C. D., Manis, P. B. & Linden, D. J. Polarity of long-term synaptic gain
841 change is related to postsynaptic spike firing at a cerebellar inhibitory synapse.
842 *Neuron* **21**, 827–835 (1998).
- 843 [71] Aizenman, C. D. & Linden, D. J. Rapid, synaptically driven increases in the
844 intrinsic excitability of cerebellar deep nuclear neurons. *Nature neuroscience* **3**,
845 109–111 (2000).
- 846 [72] Zheng, N. & Raman, I. M. Synaptic inhibition, excitation, and plasticity in
847 neurons of the cerebellar nuclei. *The Cerebellum* **9**, 56–66 (2010).
- 848 [73] Hull, C. Prediction signals in the cerebellum: beyond supervised motor learning.
849 *Elife* **9**, e54073 (2020).
- 850 [74] Kawato, M., Furukawa, K. & Suzuki, R. A hierarchical neural-network model for
851 control and learning of voluntary movement. *Biological cybernetics* **57**, 169–185
852 (1987).
- 853 [75] Wolpert, D. M., Miall, R. C. & Kawato, M. Internal models in the cerebellum.
854 *Trends in cognitive sciences* **2**, 338–347 (1998).
- 855 [76] Vijayakumar, S. & Schaal, S. Locally weighted projection regression: An o
856 (n) algorithm for incremental real time learning in high dimensional space. In
857 *Proceedings of the Seventeenth International Conference on Machine Learning*
858 (*ICML 2000*), vol. 1, 288–293 (2000).

- 860 [77] Robinson, D. Adaptive gain control of vestibuloocular reflex by the cerebellum.
861 *Journal of neurophysiology* **39**, 954–969 (1976).
- 862 [78] Miles, F. & Lisberger, S. Plasticity in the vestibulo-ocular reflex: a new hypothesis.
863 *Annual review of neuroscience* **4**, 273–299 (1981).
- 864 [79] Lac, S., Raymond, J. L., Sejnowski, T. J. & Lisberger, S. G. Learning and
865 memory in the vestibulo-ocular reflex. *Annual review of neuroscience* **18**, 409–
866 441 (1995).
- 867 [80] Ito, M. Cerebellar learning in the vestibulo-ocular reflex. *Trends in cognitive
868 sciences* **2**, 313–321 (1998).
- 869 [81] Boyden, E. S., Katoh, A. & Raymond, J. L. Cerebellum-dependent learning: the
870 role of multiple plasticity mechanisms. *Annual review of neuroscience* **27** (2004).
- 871 [82] Clopath, C., Badura, A., De Zeeuw, C. I. & Brunel, N. A cerebellar learn-
872 ing model of vestibulo-ocular reflex adaptation in wild-type and mutant mice.
873 *Journal of Neuroscience* **34**, 7203–7215 (2014).
- 874 [83] Armstrong, D. & Rawson, J. Activity patterns of cerebellar cortical neurones
875 and climbing fibre afferents in the awake cat. *The Journal of Physiology* **289**,
876 425–448 (1979).
- 877 [84] Rieubland, S., Roth, A. & Häusser, M. Structured connectivity in cerebellar
878 inhibitory networks. *Neuron* **81**, 913–929 (2014).
- 879 [85] Wang, W. X. & Lefebvre, J. L. Morphological pseudotime ordering and fate
880 mapping reveals diversification of cerebellar inhibitory interneurons. *bioRxiv*
881 (2020).
- 882 [86] Tsao, D. Y., Freiwald, W. A., Knutson, T. A., Mandeville, J. B. & Tootell, R. B.
883 Faces and objects in macaque cerebral cortex. *Nature neuroscience* **6**, 989–995
884 (2003).
- 885 [87] Zhen, Z., Fang, H. & Liu, J. The hierarchical brain network for face recognition.
886 *PloS one* **8**, e59886 (2013).
- 887 [88] Raymond, J. L. & Medina, J. F. Computational Principles of Supervised Learn-
888 ing in the Cerebellum. *Annual Review of Neuroscience* **41**, 233–253 (2018).
- 889 [89] Walter, J. T. & Khodakhah, K. The Linear Computational Algorithm of Cere-
890 bellar Purkinje Cells. *Journal of Neuroscience* **26**, 12861–12872 (2006). URL
891 <https://www.jneurosci.org/content/26/50/12861>.
- 892 [90] Jörntell, H., Bengtsson, F., Schonewille, M. & De Zeeuw, C. I. Cerebellar molec-
893 ular layer interneurons—computational properties and roles in learning. *Trends
894 in neurosciences* **33**, 524–532 (2010).
- 895 [91] Sotelo, C. Molecular layer interneurons of the cerebellum: developmental and
896 morphological aspects. *The Cerebellum* **14**, 534–556 (2015).
- 897 [92] Kim, J. & Augustine, G. J. Molecular layer interneurons: key elements of cere-
898 bellar network computation and behavior. *Neuroscience* **462**, 22–35 (2021).
- 899 [93] Palay, S. L. & Chan-Palay, V. *Cerebellar cortex: cytology and organization*
900 (Springer Science & Business Media, 2012).

- 901 [94] Pouzat, C. & Hestrin, S. Developmental regulation of basket/stellate cell→ purk-
902 inje cell synapses in the cerebellum. *Journal of Neuroscience* **17**, 9104–9112
903 (1997).
- 904 [95] Dean, P., Porrill, J., Ekerot, C.-F. & Jörntell, H. The cerebellar microcircuit
905 as an adaptive filter: experimental and computational evidence. *Nature Reviews
906 Neuroscience* **11**, 30–43 (2010).
- 907 [96] Rowan, M. J. *et al.* Graded control of climbing-fiber-mediated plasticity and
908 learning by inhibition in the cerebellum. *Neuron* **99**, 999–1015 (2018).
- 909 [97] Isope, P. & Barbour, B. Properties of unitary granule cell→ purkinje cell synapses
910 in adult rat cerebellar slices. *Journal of Neuroscience* **22**, 9668–9678 (2002).
- 911 [98] Yang, Y. & Lisberger, S. G. Purkinje-cell plasticity and cerebellar motor learning
912 are graded by complex-spike duration. *Nature* **510**, 529–532 (2014).
- 913 [99] Wagner, M., Kim, T., Savall, J., Schnitzer, M. & Luo, L. Cerebellar granule cells
914 encode the expectation of reward. *Nature* **544** (2017).
- 915 [100] Giovannucci, A. *et al.* Cerebellar granule cells acquire a widespread predictive
916 feedback signal during motor learning. *Journal of Neurophysiology* **20**, 727–734
917 (2017). URL <https://www.nature.com/articles/nn.4531>.
- 918 [101] Tremblay, R., Lee, S. & Rudy, B. Gabaergic interneurons in the neocortex: from
919 cellular properties to circuits. *Neuron* **91**, 260–292 (2016).
- 920 [102] Bradbury, J. *et al.* JAX: composable transformations of Python+NumPy pro-
921 grams (2018). URL <http://github.com/google/jax>.
- 922 [103] Babuschkin, I. *et al.* The DeepMind JAX Ecosystem (2020). URL <http://github.com/deepmind>.
- 923 [104] Ekaba, B. Google colaboratory. In: Building Machine Learning and Deep Learn-
924 ing Models on Google Cloud Platform. Apress, Berkeley, CA. (2019). URL
925 https://doi.org/10.1007/978-1-4842-4470-8_7.
- 926 [105] Kingma, D. P. & Ba, J. Adam: A method for stochastic optimization. *arXiv
927 preprint arXiv:1412.6980* (2014).
- 928 [106] Tanno, R., Arulkumaran, K., Alexander, D., Criminisi, A. & Nori, A. Adaptive
929 neural trees. In Chaudhuri, K. & Salakhutdinov, R. (eds.) *Proceedings of the
930 36th International Conference on Machine Learning*, vol. 97 of *Proceedings of
931 Machine Learning Research*, 6166–6175 (PMLR, Long Beach, California, USA,
932 2019). URL <http://proceedings.mlr.press/v97/tanno19a.html>.
- 933 [107] Arik, S. O. & Pfister, T. Tabnet: Attentive interpretable tabular learning. *Pro-
934 ceedings of the AAAI Conference on Artificial Intelligence (To Appear)* (2021).
- 935 [108] Kostadinov, D., Beau, M., Pozo, M. B. & Häusser, M. Predictive and reactive
936 reward signals conveyed by climbing fiber inputs to cerebellar purkinje cells.
937 *Nature Neuroscience* **22**, 950–962 (2019).
- 938 [109] Hippenmeyer, S. *et al.* A developmental switch in the response of drg neurons
939 to ets transcription factor signaling. *PLoS biol* **3**, e159 (2005).
- 940

- 941 [110] Dana, H. *et al.* High-performance calcium sensors for imaging activity in neuronal
942 populations and microcompartments. *Nature Methods* **16**, 649–657 (2019).
- 943 [111] Pachitariu, M. *et al.* Suite2p: beyond 10,000 neurons with standard two-photon
944 microscopy. *Biorxiv* (2017).
- 945 [112] Deneux, T. *et al.* Accurate spike estimation from noisy calcium signals for ul-
946 trafast three-dimensional imaging of large neuronal populations *in vivo*. *Nature*
947 *communications* **7**, 1–17 (2016).
- 948 [113] Storey, J. D. A direct approach to false discovery rates. *Journal of the Royal*
949 *Statistical Society: Series B (Statistical Methodology)* **64**, 479–498 (2002).
- 950 [114] LeCun, Y., Cortes, C. & Burges, C. Mnist handwritten digit database. *ATT*
951 *Labs [Online]*. Available: <http://yann.lecun.com/exdb/mnist> **2** (2010).

952 Supplementary Information

953 Difference between GLNs and DGNs

954 The main difference between the DGNs and Gated Linear Networks (GLNs) is in the
 955 interplay of the gating functions and synaptic weights. To understand this, it is helpful
 956 to write Eq. (1) in the form

$$r_{k,i} = \phi \left(\sum_{j=0}^{n_{k-1}} \left[\sum_{b=1}^{B_{k,i}} g_{k,i}^b(\mathbf{x}) w_{k,ij}^b \right] h_{k-1,j} \right). \quad (13)$$

957 The term in brackets is the effective weight, which, for each j , consists of a sum over
 958 branches. The gates, $g_{k,i}^b$, can be either zero or 1; since there are $B_{k,i}$ of them, there
 959 are $2^{B_{k,i}}$ possible effective weights. For the GLN, on the other hand, $r_{k,i}$ is given by
 960

$$r_{k,i} = \phi \left(\sum_{j=0}^{n_{k-1}} w_{k,ij}^{\beta(\mathbf{x})} h_{k-1,j} \right). \quad (14)$$

961 The difference between this and Eq. (13) is the the term in brackets has been replaced
 962 by a single weight, $w_{k,ij}^{\beta(\mathbf{x})}$. However, the index $\beta(\mathbf{x})$ can take on $2^{B_{k,i}}$ values, so there
 963 are just as many weights in the GLN as effective weights in the DGN. The value of
 964 $\beta(\mathbf{x})$ is determined by the input, and is given by the binary string (suppressing the \mathbf{x}
 965 dependence for clarity)

$$\beta = g_{k,i}^1 g_{k,i}^2 \dots g_{k,i}^{B_{k,i}}. \quad (15)$$

966 (So if there were 5 gates, and the 3rd and 5th were on while the others were off, the
 967 term on the right would be 00101, and β would be 5.)

968 In our experience, for similar B , DGNs and GLNs perform equally well. Computationally,
 969 DGNs are more memory efficient, as B weight vectors need to be stored per
 970 neuron as opposed to 2^B for GLNs. However, this comes at the cost of more operations,
 971 as there is an additional sum over branches (the term in brackets in Eq. (13)).

972 The difference in the number of parameters translates to a difference in inductive
 973 bias. GLNs are less prone to catastrophic forgetting compared to DGNs, as only one
 974 weight vector per neuron is updated for each input. This, however, means that DGNs
 975 are better than GLNs at learning new tasks – so long as there is some shared structure.

976 Convexity

977 Here we show that the loss is convex with respect to the weights in the previous layer.
 978 Temporarily dropping indices for clarity, the loss, $\ell(r^*, r)$, is given in terms of the

981 weight vector, \mathbf{w} , as $\ell(r^*, r) = \ell(r^*, \phi(h))$ with $h = \mathbf{c} \cdot \mathbf{w}$ (see Eq. (1)). If $\ell(r^*, \phi(h))$ is
 982 convex in h , then ℓ is also convex in \mathbf{w} , since h is a linear function of \mathbf{w} .

983 For quadratic loss, Eq. (5), ϕ is the identity, so $\ell(r^*, \phi(h)) = \frac{1}{2}(r^* - h)^2$. This is
 984 obviously convex in h , and so convex in \mathbf{w} . For binary cross-entropy loss, Eq. (10),
 985 $\phi(h)$ is given by a clipped sigmoid, Eq. (9). When clipped, $\phi(h) = 0$, which is convex.
 986 When not clipped, $\phi(h) = \sigma(h) = 1/(1 + e^{-h})$, for which it is easy to show that
 987 $\partial^2 \ell(r^*, \phi(h))/\partial h^2 = \sigma(h)(1 - \sigma(h)) > 0$. Thus, again ℓ is convex in h , and so also in
 988 \mathbf{w} .

989 Optimal number of branches

990 The number of dendritic branches is one of the main factors determining the model
 991 capacity of DGNs. Having too few branches results in underfitting, as the network is
 992 not flexible enough to learn the underlying function. Having too many branches, on
 993 the other hand, can result in memorization, and thus overfitting. We have generally
 994 used 10 branches per neuron except in the Inverse Kinematics experiments, where we
 995 used up to 5000 branches, as the task measures memorization not generalization.

996 In Figure S1, we show the average accuracy in the permuted MNIST task as a func-
 997 tion of the number of branches. This inverted U-shaped relationship can be observed
 998 in most tasks (data not shown). In Figure S2, we show how the MSE improves with
 999 an increased number of branches. Because the training and test data is mixed, over-
 1000 fitting is not possible, and so performance improves monotonically with the number of
 1001 branches.

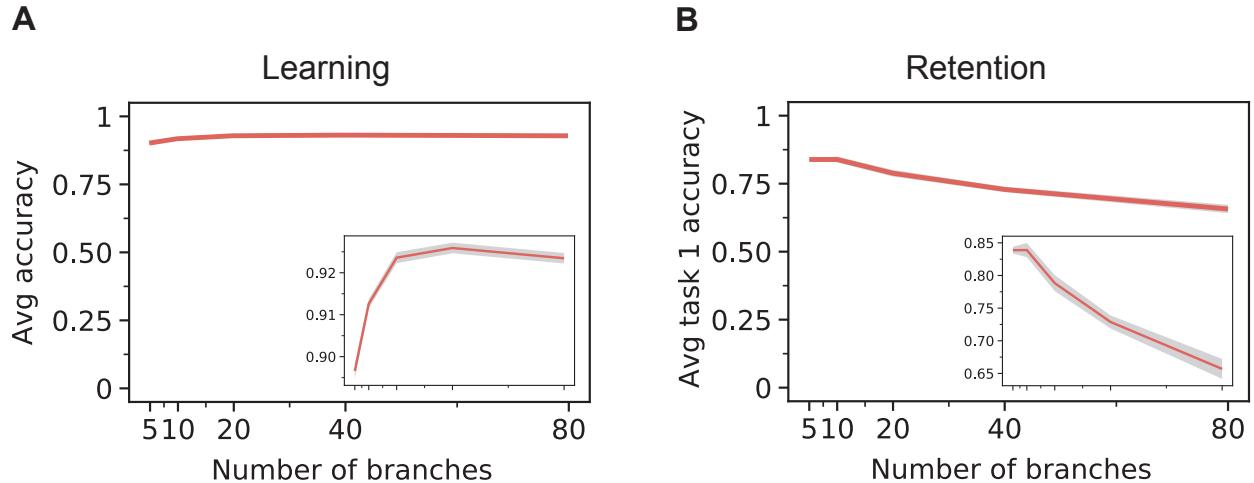


Figure S1: Permutated MNIST as a function of the number of DGN branches. **A.** Test accuracy at the end of training of each task, averaged over all 10 tasks. **B.** Test accuracy on task 1 after training on all 9 permutations. Grey areas are 99.5% confidence intervals of the results obtained from 10 models, initialised with different gating parameters and trained on differently permuted data.

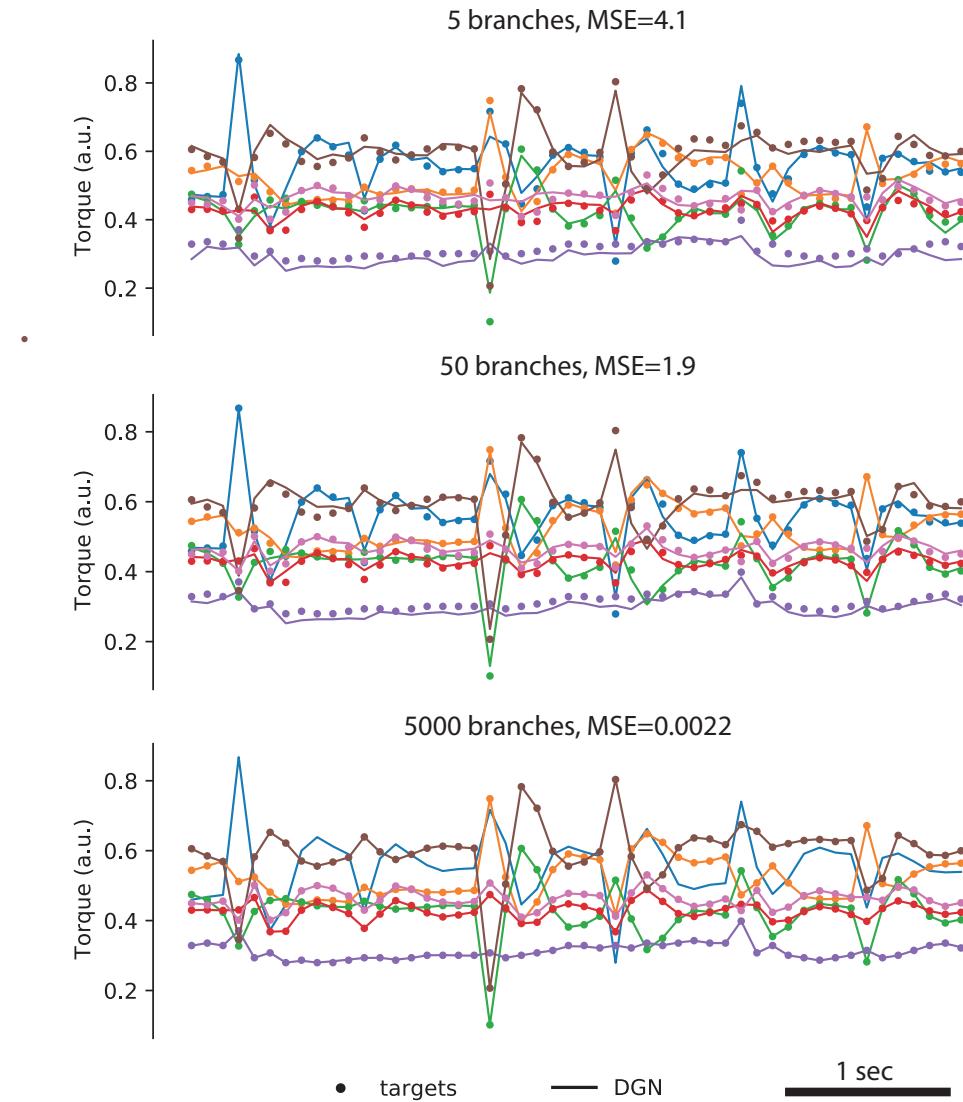


Figure S2: Sarcos solution for DGNs with 5, 50 and 5000 branches. Learning rate was 10^{-4} for 5 branches and 10^{-5} for 50 and 5,000 branches.

Inverse Kinematics

In Table S1 we compare the mean square error (MSE) obtained by DGN against baselines obtained from [31, 106, 107]. Note that, as mentioned in Methods, we (like others) used a test set that contained training examples.

Algorithm	MSE	Epochs
DGN	0.002	2000
Random forest	2.39	-
MLP	2.13	-
Stochastic decision tree	2.11	-
Gradient boosted tree	1.44	-
TabNet-S	1.25	55000
Adaptive neural tree	1.23	-
TabNet-M	0.28	55000
TabNet-L	0.14	55000
Gaussian Gated Linear Network	0.10	2000

Table S1: Test mean square error (MSE) and the number of passes over the dataset (i.e., number of epochs) for DGN with 5,000 branches versus previously published methods on the SARCOS inverse dynamics dataset [76, 106, 107]. DGN obtains the best result, by a factor of 50.

1006 Catastrophic Forgetting (permuted MNIST)

1007 **Hyperparameter selection.** We used a grid search to select the hyperparameters
 1008 for the three networks (DGN, MLP and EWC). The parameters we tested are shown
 1009 in Table S2; the ones that maximize test accuracy are in bold.

Model	learning rate	dropout	regularization const
DGN	$10^{-4}, 10^{-3}, \mathbf{10^{-2}}, 10^{-1}$	-	-
MLP	$10^{-6}, 10^{-5}, \mathbf{10^{-4}}, 10^{-3}$	Yes, No	-
EWC	$10^{-6}, 10^{-5}, \mathbf{10^{-4}}, 10^{-3}$	Yes, No	$10^2, \mathbf{10^3}, 10^4$

Table S2: For permuted MNIST, parameters swept during grid search. The best parameters (shown in bold) are the ones that maximize the average test accuracy over 20 random seeds.

1010 **Learning curves.** In Fig. S3 we show the test performance of previously learned
 1011 tasks (columns) as a function of the training across multiple tasks. To reduce clutter, a
 1012 subset of the tasks (1, 2, 4, and 8, out of 10) are shown. The top left plot (train and test
 1013 on task 1) shows that DGNs learn the first task much faster than all other methods.
 1014 The plots to the right of that show retention on task 1 while the network is sequentially
 1015 trained on subsequent tasks. MLP performances drop drastically after learning a few
 1016 new tasks, while DGN and EWC show little forgetting. This is a remarkable feat for
 1017 DGNs, which have no access to task boundaries and no explicit memory of previously
 1018 learned tasks. EWCs, on the other hand, have both. If we look at the four diagonal
 1019 plots, we see that DGN learns new tasks faster than all other methods, although the
 1020 difference gets smaller as more tasks are learned.

1021 The final accuracies across the diagonal correspond to the left panel of Fig. 4
 1022 whereas the final accuracies across the first row correspond to the right panel.

1023 VOR

1024 To obtain the smooth connectivity patterns seen in Fig. 7D, the initial weights had to
 1025 be small. Larger initial weights produced non-smooth connectivity patterns, although
 1026 the non-smoothness was different for MLPs than it was for DGNs. For MLPs, standard
 1027 Glorot intialisation led to the noisy connectivity patterns shown in Fig. S4D, bottom
 1028 panel; in contrast, to produce smooth patterns, the weights had to be scaled down by
 1029 a factor of of 100. For DGNs, scaling the initial weights up by a factor of 10 relative to
 1030 Fig. 7D produced noisy weights, but riding on a smooth background (Fig. S4D, top).

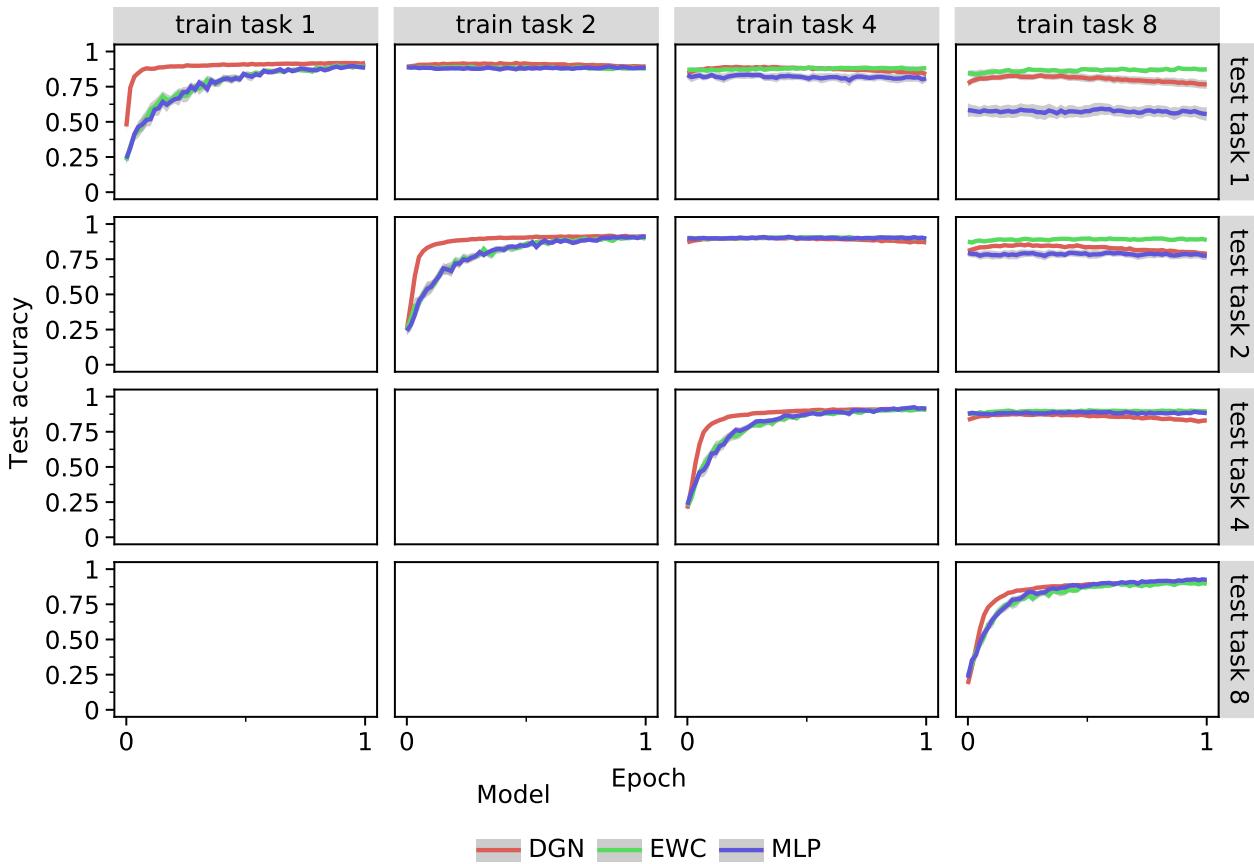


Figure S3: Retention results for permuted MNIST. Models are trained sequentially on 10 tasks, a subset of which is shown (tasks 1, 2, 4 and 8). Each column corresponds to a different stage of training (see labels on top), and each row reports test accuracy for a specific task. For example, the top row indicates performance on task 1 after being trained sequentially on tasks 1, 2, 4 and 8. Each model trains for one epoch per task; i.e., the 60,000 training examples per task are used only once. Error bars, indicated by the thickness of the lines, denote 95% confidence levels over 20 random seeds.

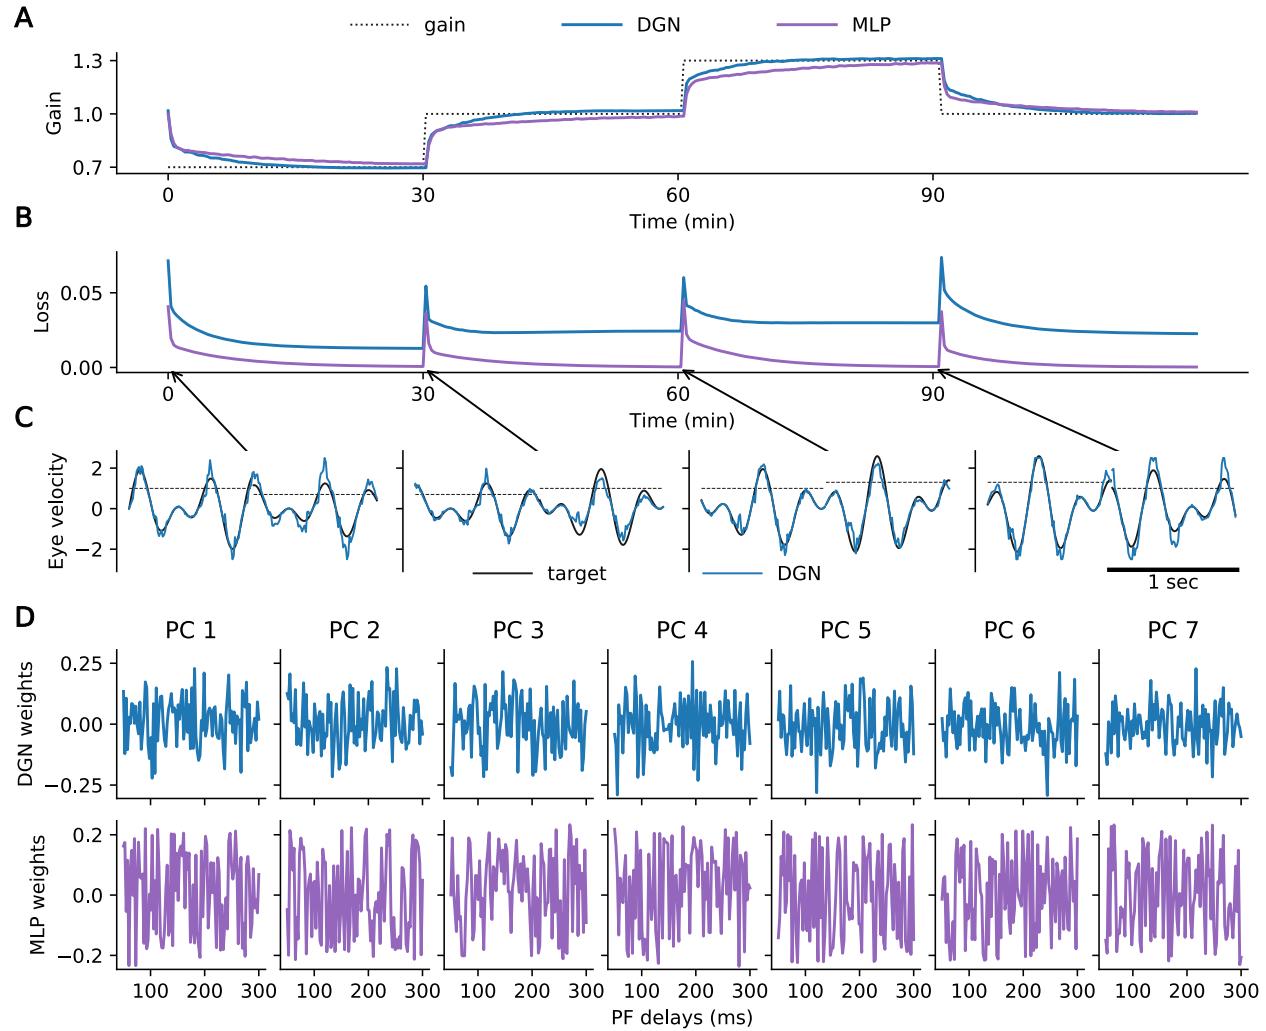


Figure S4: VOR network initialized with large, noisy weights. Same as Fig. 7, except that the training starts from large, noisy weights. For clarity, only five branches are shown in the top panel of **D** (compared to 10 in Fig. 7).

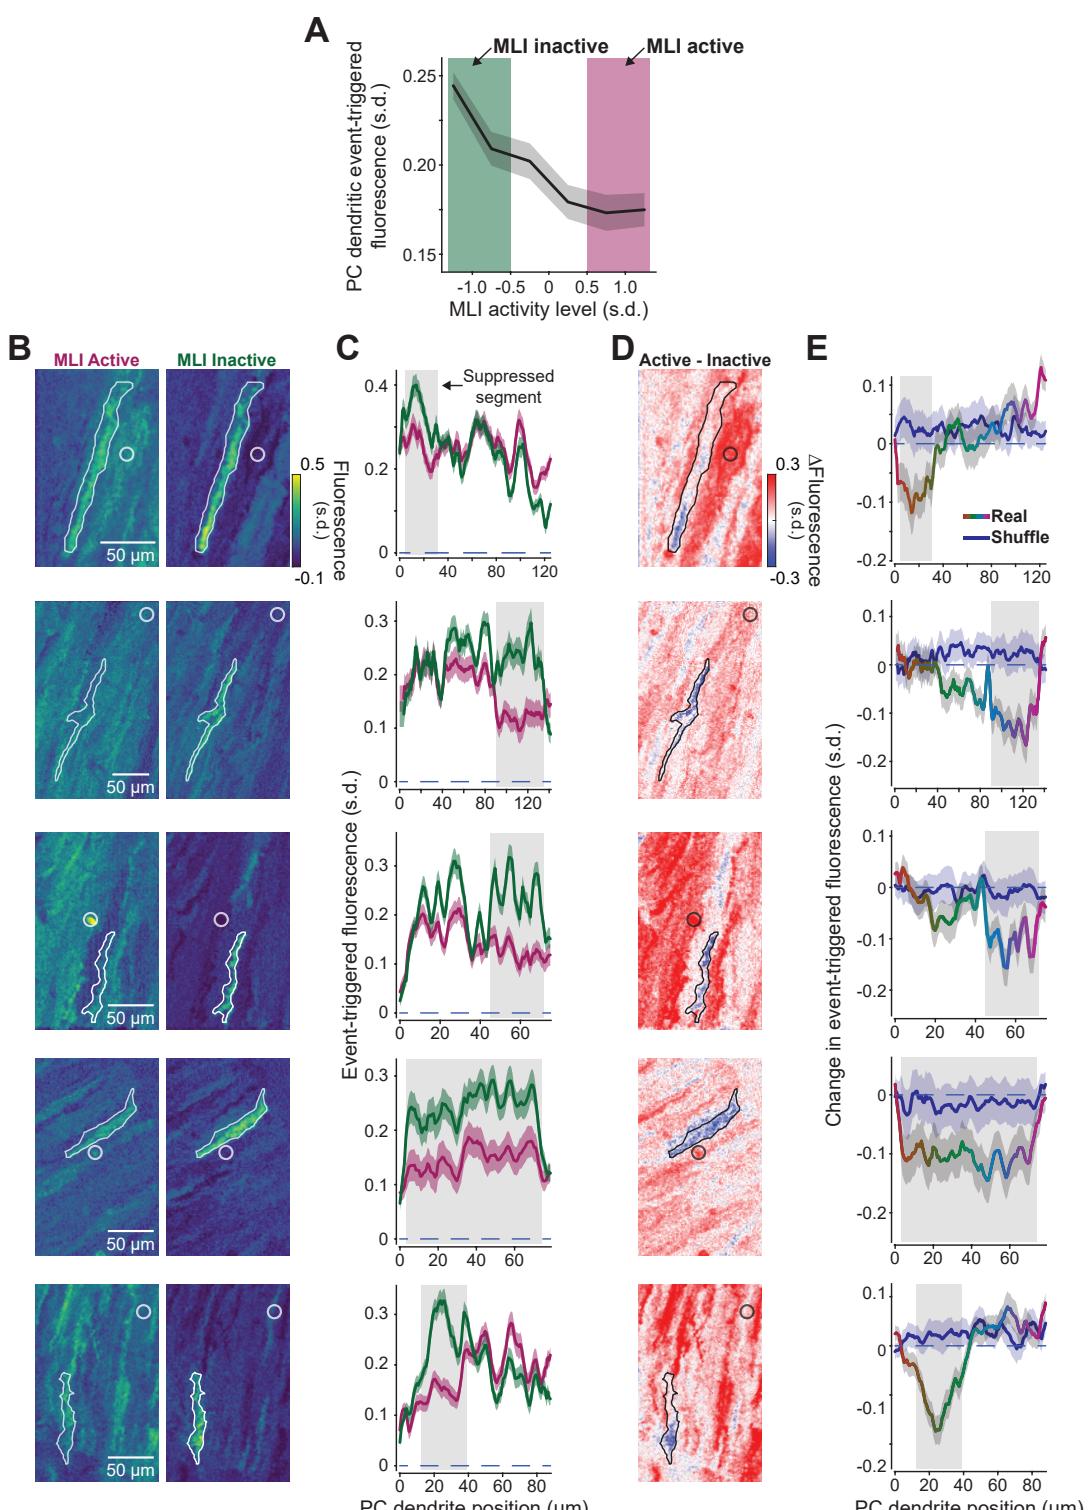


Figure S5: Suppression of Purkinje cell dendritic segments by MLIs. **A.** Event-triggered fluorescence in modulated Purkinje cell dendritic segments as a function of MLI activity level. Activity levels used for analysis in main figure are highlighted. **B.** Five additional examples of spatial event-triggered map of area surrounding Purkinje cell dendrites (contoured region of interest) when a nearby MLI (white circle;³⁹ sometimes projected from different plane) was in an active state or inactive state. **C.** Spatial profile of event-triggered fluorescence of PC dendrites shown in panel A. **D.** Difference heatmap image of event-triggered fluorescence of same dendrites shown in panels A-B. **E.** Spatial profile trace (rainbow) and shuffled trace (blue) of event-triggered difference heatmap images.

1031

Pseudocode

Algorithm S1 DGN for quadratic loss

```

1: Input: network architecture: number of layers  $K \in \mathbb{N}$ ,
   number of neurons in layer  $k$   $\{n_k \in \mathbb{N}\}$ ,
   number of branches per neuron  $i$  in layer  $k$   $\{B_{k,i} \in \mathbb{N}\}$ 
2: Input: weights  $\{w_{k,ij}^b \in \mathbb{R}\}$ 
3: Input: gating parameters  $\{v_{k,ij}^b \in \mathbb{R}\}$ ,  $\{\theta_{k,i}^b \in \mathbb{R}\}$ 
4: Input: input  $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$ 
5: Input: target  $r^* \in \mathbb{R}$ 
6: Input: learning rate  $\eta \in (0, 1)$ 
7: Input: update  $\in \{\text{TRUE}, \text{FALSE}\}$  (enables learning)
8: Output: Target prediction  $\hat{r} = r_{K,1}$  (output of neuron in last layer  $K$ )
9:  $r_{0,0} \leftarrow 1$ ;  $n_0 \leftarrow n$ ;  $r_{0,i} = x_i$  for  $i \in \{1, \dots, n\}$ 
10: for  $k \in \{1, \dots, K\}$  do {over layers}
11:    $r_{k,0} \leftarrow 1$  {bias}
12:   for  $i \in \{1, \dots, n_k\}$  do {over neurons}
13:     for  $b \in \{1, \dots, B_{k,i}\}$  do {over branches}
14:        $g_{k,i}^b \leftarrow \Theta(\sum_{j=0}^{n_{k-1}} v_{k,ij}^b x_j - \theta_{k,i}^b)$ 
15:        $r_{k,i} \leftarrow \sum_{b=1}^{B_{k,i}} g_{k,i}^b \sum_{j=0}^{n_{k-1}} w_{k,ij}^b r_{k-1,j}$ 
16:       if update then
17:         for  $b \in \{1, \dots, B_{k,i}\}$  do {over branches}
18:           if  $g_{k,i}^b > 0$  then
19:             for  $j \in \{1, \dots, n_{k-1}\}$  do {over neurons in previous layer}
20:                $w_{k,ij}^b \leftarrow w_{k,ij}^b - \eta (r_{k,i} - r^*) w_{k,ij}^b r_{k-1,j}$ 
21: return  $r_{K,1}$ 
```

1032

Here $\Theta(\cdot)$ is the Heaviside step function ($\Theta(z) = 1$ for $z > 0$ and $\Theta(z) = 0$ otherwise).

Algorithm S2 DGN for Bernoulli data

```

1: Input: network architecture: number of layers  $K \in \mathbb{N}$ ,  

    number of neurons in layer  $k$   $\{n_k \in \mathbb{N}\}$ ,  

    number of branches per neuron  $i$  in layer  $k$   $\{B_{k,i} \in \mathbb{N}\}$   

2: Input: weights  $\{w_{k,ij}^b \in \mathbb{R}\}$   

3: Input: gating parameters  $\{v_{k,ij}^b \in \mathbb{R}\}$ ,  $\{\theta_{k,i}^b \in \mathbb{R}\}$   

4: Input: precision  $\epsilon \in (0, 0.5)$   

5: Input: input  $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$   

6: Input: target  $r^* \in \{0, 1\}$   

7: Input: learning rate  $\eta \in (0, 1)$   

8: Input: update  $\in \{\text{TRUE}, \text{FALSE}\}$  (enables learning)  

9: Output: Target prediction  $\hat{r} = r_{K,1}$  (output of neuron in last layer  $K$ )  

10:  $r_{0,0} \leftarrow \sigma(1)$ ;  $n_0 \leftarrow n$ ;  $r_{0,i} = \text{CLIP}_\epsilon^{1-\epsilon}(\sigma(x_i))$  for  $i \in \{1, \dots, n\}$   

11: for  $k \in \{1, \dots, K\}$  do {over layers}  

12:    $r_{k,0} \leftarrow \sigma(1)$  {bias}  

13:   for  $j \in \{1, \dots, n_{k-1}\}$  do {over neurons in layer below}  

14:      $h_{k-1,j} \leftarrow \sigma^{-1}(r_{k-1,j})$   

15:     for  $i \in \{1, \dots, n_k\}$  do {over neurons}  

16:       for  $b \in \{1, \dots, B_{k,i}\}$  do {over branches}  

17:          $g_{k,i}^b \leftarrow \Theta(\sum_{j=0}^{n_{k-1}} v_{k,ij}^b x_j - \theta_{k,i}^b)$   

18:          $h_{k,i} \leftarrow \sum_{b=1}^{B_{k,i}} g_{k,i}^b \sum_{j=0}^{n_{k-1}} w_{k,ij}^b h_{k-1,j}$   

19:          $r_{k,i} \leftarrow \text{CLIP}_\epsilon^{1-\epsilon} \sigma(h_{k,i})$   

20:         if update then  

21:           for  $b \in \{1, \dots, B_{k,i}\}$  do {over branches}  

22:             if  $|r^* - \sigma(h_{k,i})| > \varepsilon$  then  

23:               for  $j \in \{1, \dots, n_{k-1}\}$  do {over neurons in previous layer}  

24:                  $w_{k,ij}^b \leftarrow w_{k,ij}^b - \eta(r_{k,i} - r^*) h_{k-1,j}$   

25: return  $r_{K,1}$ 

```

1033 Here, as above, $\text{CLIP}_a^b(\cdot)$ clips values between a and b ,

$$\text{CLIP}_a^b(y) \equiv \begin{cases} a & y < a \\ y & a < y < b \\ b & b \leq y \end{cases}. \quad (16)$$

1034 Also as above, $\sigma(\cdot)$ is the sigmoid function, $\sigma(z) = e^z/(1 + e^z)$. Its inverse is given by
 1035 $\sigma^{-1}(y) = \log(y/(1 - y))$.